


COMP 430/533

Intro. to Database Systems

SQL 1

Have you used SQL before?

- A. Not at all
- B. Seen it, never wrote any
- C. Written simple queries
- D. Comfortable with the basics
- E. Used it a lot




Basic concepts

Database Management System (DBMS)

Software for managing multiple databases – e.g., PostgreSQL

- Server that communicates with clients
- DB ops controlled by client, typically using SQL
- Manages users & access control
- Often given exclusive control over storage
 - Potentially no traditional file system
 - May replicate & distribute data over multiple disks or servers



We'll use three DB clients

Jupyter:

- SQL embedded within Jupyter notebooks

psql:


- Command-line
- .sql code file

user app:

- Called from main application language (Python, Java, ...)
- E.g., something like `SQLexecute("...SQL CODE...")`

Table=Relation


Heading:
Names of each
column.



Product


p_name	price	manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$39.99	GizmoWorks
Widget	\$19.99	WidgetsRUs
HyperWidget	\$203.99	Hyper

Row=tuple=record:
A single entry in the
table.
#rows = cardinality



Column=attribute=field:
A typed data entry in each row.
#columns = arity=degree

Cell=field:
One column of one row



Table=Relation

Product

p_name	price	manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$39.99	GizmoWorks
Widget	\$19.99	WidgetsRUs
HyperWidget	\$203.99	Hyper

Table is a *multiset* of rows. Not ordered. Duplicate rows allowed.

Schema defines table attributes

Product

p_name	price	manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$39.99	GizmoWorks
Widget	\$19.99	WidgetsRUs
HyperWidget	\$203.99	Hyper

Product (p_name: string, price: float, manufacturer: string)

Creating a table

Product (p_name: string, price: money, manufacturer: string)

```
CREATE TABLE Product (  
  p_name VARCHAR(50),  
  price MONEY,  
  manufacturer VARCHAR(50)  
);
```

Some SQLs:
CURRENCY,
NUMERIC(6,2)

```
CREATE TYPE ProductType (  
  p_name VARCHAR(50),  
  price MONEY,  
  manufacturer VARCHAR(50)  
);  
CREATE TABLE Product OF ProductType;
```

SQL attribute types

All types are atomic!

Traditionally. However, sets and arrays are allowed in some SQL versions. We won't discuss these.

Mostly standard:

- CHAR(n), VARCHAR(n), NCHAR(n), NVARCHAR(n)
- INT, BIGINT, SMALLINT, FLOAT, DECIMAL(m,n)
- TIMESTAMP, DATE, TIME, INTERVAL
- ...

Less standard:

- MONEY, BOOLEAN, UUID
- ...

NULL data

Any cell can be NULL.

- Signals missing value.

Possible interpretations:

- Value doesn't exist.
- Value exists but unknown.
- Value not applicable to this record.

Keys

Key = minimal set of attributes acting as a unique tuple identifier.

- Consider the universe of all potential relation data, not just what the table currently contains.

Product


p_name	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper

What is key?

- A. dept
- B. dept, number
- ✓ C. dept, number, section
- D. dept, number, section, instructor

Course

dept	number	section	instructor
MATH	101	1	Jones
MATH	101	2	Smith
MATH	102	1	Williams
PHYS	101	1	Baker
PHYS	102	2	Baker




What is key?

- A. s_name
- B. abbrev
- C. The pair s_name, abbrev
- ✓ D. Each of s_name and abbrev

State

s_name	abbrev	year_admitted
Alabama	AL	1819
Alaska	AK	1959
Arizona	AZ	1912
Arkansas	AR	1836
California	CA	1850



Primary keys

Every table should have one *primary key*.

- Each tuple must have distinct non-NULL values of primary key attributes.
- Guarantees table is a mathematical relation.

Product

p_name	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper

Product (p_name: string, price: float, manufacturer: string)

Primary keys are a DB-checked constraint

Adding a new record with duplicate or NULL primary key will fail.

- Failure reported as exception or erroneous return value.

Product

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper
Widget	15.99	GizmoWorks

Product (p_name:
string, price: float,
manufacturer: string)

We will see other kinds
of DB-checked
constraints.

What data should the primary key be?

Natural: Data field(s) needed in table anyways.

Synthetic: “Made-up” just for the purpose of being an identifier.

PostgreSQL: UUID type.

Pros/cons discussed later in design section.

Creating a table with primary key

Product (p_name: string, price: money, manufacturer: string)

```
CREATE TABLE Product (  
    p_name VARCHAR(50),  
    price MONEY,  
    manufacturer VARCHAR(50),  
    PRIMARY KEY (p_name)  
);
```

```
CREATE TABLE Product (  
    p_name VARCHAR(50) PRIMARY KEY,  
    price MONEY,  
    manufacturer VARCHAR(50)  
);
```

NOT NULL constraint


Course (c_id: UUID, c_name: string, instructor: string)

```
CREATE TABLE Course (  
  c_id UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
  c_name VARCHAR(50) NOT NULL,  
  instructor VARCHAR(50)  
);
```

Adding records to a table

```
INSERT INTO Product VALUES  
('MiniGizmo', 15.99, 'GizmoWoks'),  
('MiniWidget', 21.99, 'WidgetsRUs');  
INSERT INTO Product (p_name, manufacturer) VALUES  
('NanoWidget', 'WidgetsRUs');
```

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper



<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper
MiniGizmo	15.99	GizmoWorks
MiniWidget	21.99	WidgetsRUs
NanoWidget	NULL	WidgetsRUs

Activity: Creating a table

02a-table.ipynb

```
CREATE TABLE Course (...);  
INSERT INTO Course ...;
```

Course

<u>dept</u>	<u>number</u>	<u>section</u>	instructor
MATH	101	1	Jones
MATH	101	2	Smith
MATH	102	1	Williams
PHYS	101	1	Baker
PHYS	102	2	Baker

Course (dept, number, section, instructor)

Activity partial solution

Course (dept, number, section, instructor)

```
CREATE TABLE Course (  
    dept CHAR(4),  
    number CHAR(3),  
    section INT DEFAULT 1,  
    instructor VARCHAR(50),  
    PRIMARY KEY (dept, number, section)  
);
```

A convenient option.

Simple queries

SELECT-FROM-WHERE

```
SELECT attributes  
FROM tables  
WHERE conditions;
```


Results in a new table, which can be returned or stored.

Selecting some records

```
SELECT *  
FROM Product  
WHERE Price > 20;
```

Product

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper



<u>p_name</u>	price	manufacturer
Powergizmo	39.99	GizmoWorks
HyperWidget	203.99	Hyper

More selection examples

```
SELECT *  
FROM Product  
WHERE price IS NOT NULL;
```

```
SELECT *  
FROM Product  
WHERE price BETWEEN 20 AND 40;
```



Inclusive

```
SELECT *  
FROM Product  
WHERE price > 20 AND manufacturer = 'GizmoWorks';
```

```
SELECT *  
FROM Product  
WHERE manufacturer IN ('GizmoWorks', 'WidgetsRUs');
```

Selection with pattern matching

```
SELECT *  
FROM Product  
WHERE p_name LIKE '%Gizmo%';
```


- % = Match any sequence of 0-or-more characters
- _ = Match any single character
- [abc] = Match any one character listed
- [a-c] = Match any one character in range

Projecting some fields

```
SELECT  p_name, manufacturer
FROM    Product;
```

Product

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper



<u>p_name</u>	manufacturer
Gizmo	GizmoWorks
Powergizmo	GizmoWorks
Widget	WidgetsRUs
HyperWidget	Hyper


Answer (p_name, manufacturer)

Combining selection & projection

```
SELECT p_name, manufacturer  
FROM Product  
WHERE price > 20;
```

Product

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper



<u>p_name</u>	manufacturer
Powergizmo	GizmoWorks
HyperWidget	Hyper

Are results necessarily distinct?

Product


<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper

```
SELECT manufacturer
FROM Product;
```

- A. Yes
- ✓ B. No

manufacturer
GizmoWorks
WidgetsRUs
Hyper

manufacturer
GizmoWorks
GizmoWorks
WidgetsRUs
Hyper




Making results distinct

```
SELECT DISTINCT manufacturer  
FROM Product;
```

Product

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper



manufacturer
GizmoWorks
WidgetsRUs
Hyper

Ensures results are a set.

Query semantics – set notation

```
SELECT [DISTINCT] a1, a2, ..., am  
FROM T  
WHERE FilterCondition;
```

$\{(a_1, a_2, \dots, a_m) \mid \text{FilterCondition}\}$ — Multisets by default.
Sets with DISTINCT.

Computation in SELECT clauses

```
SELECT location, time, celsius * 1.8 + 32 AS fahrenheit  
FROM SensorReading;
```

```
SELECT player_id, Floor(height) AS feet, (height – Floor(height)) * 12 AS inches  
FROM Player;
```

```
SELECT student_id, CASE WHEN lastname < 'N' THEN 1 ELSE 2 END AS group  
FROM Student;
```


Use AS. Without it, SQL will create a default field name.

Sorting

```
SELECT p_name, manufacturer  
FROM Product  
ORDER BY price DESC, manufacturer;
```

Product

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper



<u>p_name</u>	manufacturer
HyperWidget	Hyper
Powergizmo	GizmoWorks
Gizmo	GizmoWorks
Widget	WidgetsRUs

Multiset semantics vs. Sorting

Table rows are unordered, except when they're ordered.



More accurately, unless you use ORDER BY:

- Can't assume anything about ordering.
- Ordering depends on implementation, which can vary.
- Queries don't necessarily maintain order of original table.

Subset of results


Some SQLs:
SELECT TOP 2 ...

```
SELECT p_name, manufacturer
FROM Product
LIMIT 2;
```

Takes first N rows.
Usually used with
ORDER BY.

Product

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper



<u>p_name</u>	manufacturer
Gizmo	GizmoWorks
Powergizmo	GizmoWorks

Activity: Writing queries

02b-queries.ipynb

Some details

NULL semantics

NULL is not a value. It is the lack of a value.

In numeric operations:

- $f(\text{NULL}) \rightarrow \text{NULL}$

In Boolean operations, we use 3-value logic (FALSE, UNKNOWN, TRUE):

- $\text{NULL} = \text{'Houston'} \rightarrow \text{UNKNOWN}$
- $\text{NULL} = \text{NULL} \rightarrow \text{UNKNOWN}$

A	NOT A
F	T
U	U
T	F

A AND B		B		
		F	U	T
A	F	F	F	F
	U	F	U	U
	T	F	U	T

A OR B		B		
		F	U	T
A	F	F	U	T
	U	U	U	T
	T	T	T	T

Syntax details

- Strings use single quotes, not double
- Equality test uses single =, not double
- Amount of whitespace doesn't matter

'Houston'

x = 5

Syntax details – case sensitivity

- Case insensitive
 - Keywords: SELECT, NULL
 - Table names: Product
 - Function names: Count()
- Depends on SQL version
 - Attribute names: product_id
- Case sensitive
 - String literals: 'HOUSTON', 'Houston', 'houston'

Some useful online SQL references

Good for quick checks of statement syntax and options.

- www.w3schools.com/sql/
- www.tutorialspoint.com/sql/
- www.sqlcommands.net/
- www.techonthenet.com/sql/

Each flavor of DBMS has its own reference.

Using multiple tables

The need for multiple tables

Using a single table leads to repeating data.

- Provides the opportunity for inconsistency
- Requires more storage space & I/O time

Product

<u>p_name</u>	price	manufacturer	address	city	state
Gizmo	19.99	GizmoWorks	123 Gizmo St.	Houston	TX
Powergizmo	39.99	GizmoWorks	123 Gizmo St.	Houston	TX
Widget	19.99	WidgetsRUs	20 Main St.	New York	NY
HyperWidget	203.99	Hyper	1 Mission Dr.	San Francisco	CA

Product (p_name, price, manufacturer, address, city, state)

Using multiple tables

Product

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper

Product (p_name, price, manufacturer)

Company

<u>c_name</u>	address	city	state
GizmoWorks	123 Gizmo St.	Houston	TX
WidgetsRUs	20 Main St.	New York	NY
Hyper	1 Mission Dr.	San Francisco	CA

Company (c_name, address, city, state)

Later in course: Deciding what fields belong in what tables.

Foreign keys & referential integrity

Product

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper

Company

<u>c_name</u>	address	city	state
GizmoWorks	123 Gizmo St.	Houston	TX
WidgetsRUs	20 Main St.	New York	NY
Hyper	1 Mission Dr.	San Francisco	CA

Product's manufacturer is a *foreign key*.

- Foreign keys always refer to primary keys.
- We want to enforce *referential integrity* – each manufacturer is a c_name.

Creating a table with a foreign key

```
CREATE TABLE Product (  
    p_name VARCHAR(50),  
    price MONEY,  
    manufacturer VARCHAR(50),  
    PRIMARY KEY (p_name),  
    FOREIGN KEY (manufacturer) REFERENCES Company (c_name )  
);
```

```
CREATE TABLE Company (  
    c_name VARCHAR(50),  
    address VARCHAR(50),  
    city VARCHAR(50),  
    state CHAR(2),  
    PRIMARY KEY (c_name)  
);
```

Foreign key represents a dependence

Product

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
...

Company

<u>c_name</u>	address	city	state
GizmoWorks	123 Gizmo St.	Houston	TX
...

Product conceptually dependent on Company to elaborate details.
Product data dependent on Company data being entered.

```
CREATE TABLE Product (  
    ....  
    FOREIGN KEY (manufacturer) REFERENCES Company (c_name)  
);
```

```
CREATE TABLE Company (...);
```

Product def'n dependent on Company def'n.

Foreign keys vs. pointers

Foreign keys relate tables, or equivalently, sets of attributes.

- Repeats data to connect individual records.
- One relationship between tables vs. many pointers between table records.

Problems with data pointers:


- Difficult to maintain pointers when data changes, moves, or gets copied, esp. with concurrency.
- We want to avoid making queries be directionally-biased.

In what ways do we relate tables?


Explore design issues & common patterns later.

Two main building blocks:

- One-to-many relationships



- Many-to-many relationships



One-to-many / many-to-one relationships

Company

<u>c_name</u>	address	city	state
GizmoWorks	123 Gizmo St.	Houston	TX
WidgetsRUs	20 Main St.	New York	NY
Hyper	1 Mission Dr.	San Francisco	CA

Company (c_name, address, city, state)

Product

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper

Product (p_name, price, manufacturer)

Each Company can have **many** Products.

Each Product is made by exactly **one** Company.

Lookup tables – common use of one-to-many

Lookup table

Ref. integrity
ensures
Company's
state is in list.

State

<u>state</u>
AL
...
TX
WV

State (state)

Company

<u>c_name</u>	address	city	state
GizmoWorks	123 Gizmo St.	Houston	TX
WidgetsRUs	20 Main St.	New York	NY
Hyper	1 Mission Dr.	San Francisco	CA

Company (c_name, address, city, state)

Each State can have **many** Companies.

Each Company is in **one** State.

Many-to-many relationships

Student

s_id	first_name	last_name
S01	John	Smith
S02	Mary	Wallace
S03	Sue	Roper
S04	Mark	Jones

Enrollment

s_id	crn	grade
S01	01234	A
S01	46117	C
S02	01234	B

Course

crn	dept	number
01234	COMP	140
15134	COMP	160
46117	ELEC	220

Each Student can be enrolled in **many** Courses.


Each Course has **many** Students.

What is Student's primary key?

- ✓ A. s_id
- B. s_id, first_name, last_name
- C. first_name, Last_name
- D. last_name

Student

s_id	first_name	last_name
S01	John	Smith
S02	Mary	Wallace
S03	Sue	Roper
S04	Mark	Jones




What is Course's primary key?

- ✓ A. crn
- B. crn, dept, number
- C. dept, number
- D. number

Course

crn	dept	number
01234	COMP	140
15134	COMP	160
46117	ELEC	220




What is Enrollment’s primary key?

- A. s_id
- ✓B. s_id, crn
- C. s_id, crn, grade
- D. crn
- E. grade


Enrollment

s_id	crn	grade
S01	01234	A
S01	46117	C
S02	01234	B



What is a foreign key in Student?

- A. s_id
- B. first_name, last_name
- ✓ C. N/A



Student (s_id, first_name, last_name)

Enrollment (s_id, crn, grade)


Course (crn, dept, number)

What is a foreign key in Course?

A. crn

B. dept, number

✓ C. N/A




Student (s_id, first_name, last_name)

Enrollment (s_id, crn, grade)

Course (crn, dept, number)

What is a foreign key in Enrollment?

- A. s_id
- B. crn
- C. The pair sid, crn
- ✓ D. Each of s_id and crn



Student (s_id, first_name, last_name)

Enrollment (s_id, crn, grade)

Course (crn, dept, number)

Creating a table with a foreign key

```
CREATE TABLE Student (  
  s_id CHAR(10),  
  first_name VARCHAR(50),  
  last_name VARCHAR(50),  
  PRIMARY KEY (s_id)  
);
```

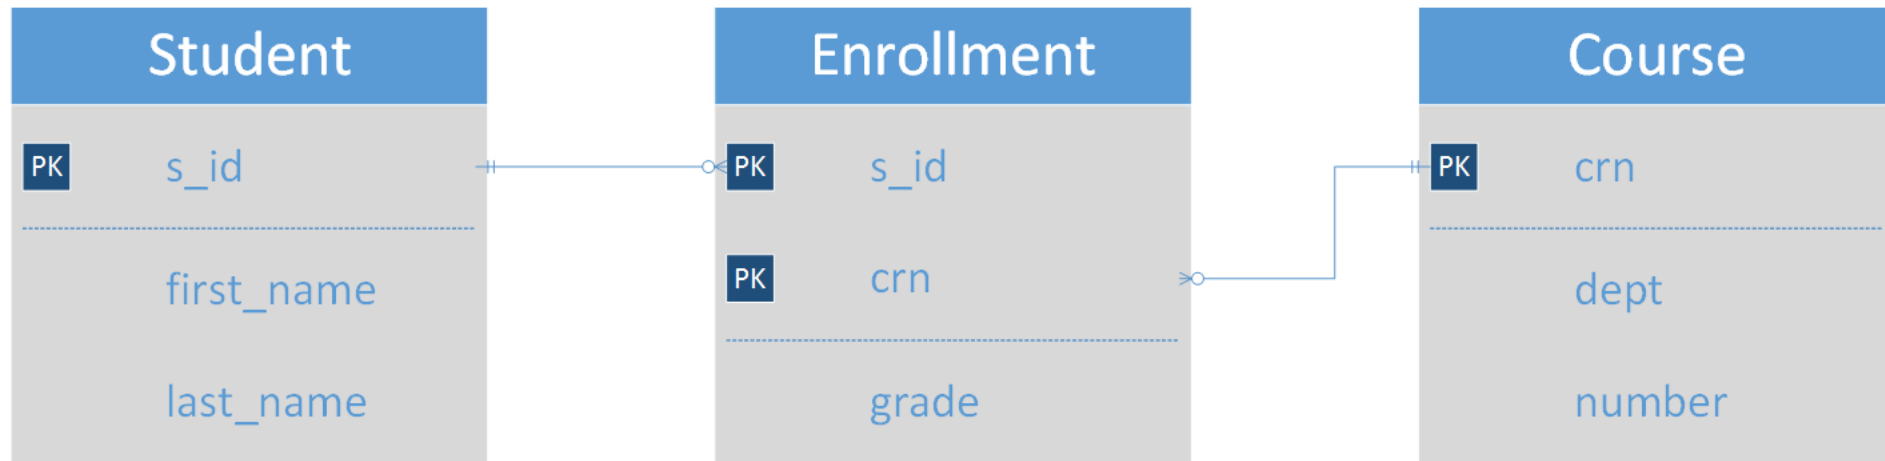
```
CREATE TABLE Course (  
  crn CHAR(10),  
  dept CHAR(4),  
  number CHAR(3),  
  PRIMARY KEY (crn)  
);
```

```
CREATE TABLE Enrollment (  
  s_id CHAR(10),  
  crn CHAR(10),  
  grade CHAR(2),  
  PRIMARY KEY (s_id, crn),  
  FOREIGN KEY (s_id) REFERENCES Student (s_id),  
  FOREIGN KEY (crn) REFERENCES Course (crn)  
);
```

Another representation – ER diagram

Style: Crow's foot

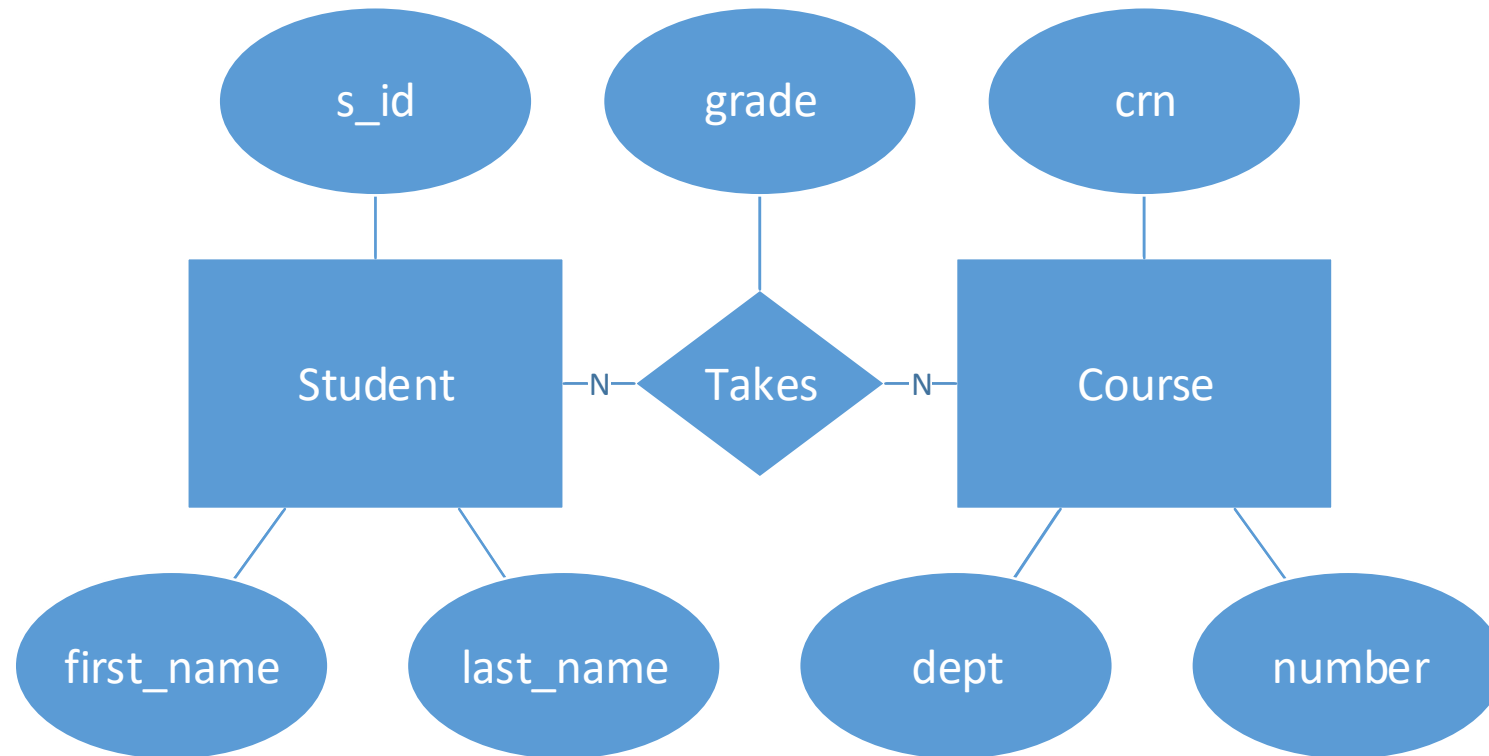
Conceptual level: Physical



A more abstract ER diagram

Style: Chen

Conceptual level: Logical



Joining tables – inner joins


Product

<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper

Company

<u>c_name</u>	address	city	state
GizmoWorks	123 Gizmo St.	Houston	TX
WidgetsRUs	20 Main St.	New York	NY
Hyper	1 Mission Dr.	San Francisco	CA

By far, the most common kind of join. See other kinds later.



<u>p_name</u>	price	manufacturer	<u>c_name</u>	address	city	state
Gizmo	19.99	GizmoWorks	GizmoWorks	123 Gizmo St.	Houston	TX
Powergizmo	39.99	GizmoWorks	GizmoWorks	123 Gizmo St.	Houston	TX
Widget	19.99	WidgetsRUs	WidgetsRUs	20 Main St.	New York	NY
HyperWidget	203.99	Hyper	Hyper	1 Mission Dr.	San Francisco	CA

Joining tables – INNER JOIN syntax

Product

<u>p_name</u>	price	manufacturer
---------------	-------	--------------

Company

<u>c_name</u>	address	city	state
---------------	---------	------	-------

```
SELECT *  
FROM Product  
INNER JOIN Company ON manufacturer = c_name;
```

```
SELECT *  
FROM Company  
INNER JOIN Product ON manufacturer = c_name;
```

```
SELECT *  
FROM Product, Company  
WHERE manufacturer = c_name;
```

Old style – deprecated & error-prone



<u>p_name</u>	price	manufacturer	<u>c_name</u>	address	city	state
Gizmo	19.99	GizmoWorks	GizmoWorks	123 Gizmo St.	Houston	TX
Powergizmo	39.99	GizmoWorks	GizmoWorks	123 Gizmo St.	Houston	TX
Widget	19.99	WidgetsRUs	WidgetsRUs	20 Main St.	New York	NY
HyperWidget	203.99	Hyper	Hyper	1 Mission Dr.	San Francisco	CA

Resolving attribute name conflicts

Person (name, address, works_for)

Company (name, address)

```
SELECT Person.name, Person.address  
FROM Person  
INNER JOIN Company ON Person.works_for = Company.name;
```

```
SELECT p.name, p.address  
FROM Person p  
INNER JOIN Company c ON p.works_for = c.name;
```

Selection & Projection

```
SELECT p_name  
FROM Product  
INNER JOIN Company ON manufacturer = c_name  
WHERE state = 'TX';
```

Join condition

Filter condition

1) Join

p_name	price	manufacturer	c_name	address	city	state
Gizmo	19.99	GizmoWorks	GizmoWorks	123 Gizmo St.	Houston	TX
Powergizmo	39.99	GizmoWorks	GizmoWorks	123 Gizmo St.	Houston	TX
Widget	19.99	WidgetsRUs	WidgetsRUs	20 Main St.	New York	NY
HyperWidget	203.99	Hyper	Hyper	1 Mission Dr.	San Francisco	CA

Selection & Projection

```
SELECT p_name  
FROM Product  
INNER JOIN Company ON manufacturer = c_name  
WHERE state = 'TX';
```

Join condition

Filter condition

2) Selection

p_name	price	manufacturer	c_name	address	city	state
Gizmo	19.99	GizmoWorks	GizmoWorks	123 Gizmo St.	Houston	TX
Powergizmo	39.99	GizmoWorks	GizmoWorks	123 Gizmo St.	Houston	TX

Selection & Projection

```
SELECT p_name  
FROM Product  
INNER JOIN Company ON manufacturer = c_name  
WHERE state = 'TX';
```

Join condition

Filter condition

3) Projection

p_name
Gizmo
Powergizmo

Semantics – set notation

```
SELECT [DISTINCT] T0.a1, T0.a2, ..., Tn.am  
FROM T0  
INNER JOIN T1 ON JoinCondition1  
...  
INNER JOIN Tn ON JoinConditionn  
WHERE FilterCondition;
```

$$\{(T_0.a_1, T_0.a_2, \dots, T_n.a_m) \mid$$
$$\text{JoinCondition}_1 \text{ AND}$$
$$\dots$$
$$\text{JoinCondition}_n \text{ AND}$$
$$\text{FilterCondition}\}$$

Multisets by default.
Sets with DISTINCT.

Activity: Writing multi-table queries

02c-queries.ipynb

Add, modify, delete data

Data not static

Databases are built to be used.

Forgetting this is a common student mistake.



Adding data

```
INSERT INTO Student VALUES  
('S123456789', 'John', 'Doe', 'Houston', 'TX');
```

```
INSERT INTO Student VALUES  
('S111111111', 'Mary', 'Jones', 'Houston', 'TX'),  
('S222222222', 'Joe", 'Wallis', 'Kansas City', 'KS');
```

```
INSERT INTO Student (id, first_name, last_name) VALUES  
('S987654321', 'Jane', 'Smith');
```

Saving query results into a table

```
INSERT INTO Student  
SELECT ...;
```

```
INSERT INTO Student (id, first_name, last_name)  
SELECT ...;
```

Existing table

```
SELECT ...  
INTO    new_table_name  
...;
```

New table

Modifying & deleting data

```
UPDATE Student  
SET city = 'Austin', state = 'TX'  
WHERE id = 'S123456789';
```

```
DELETE FROM Student;  
  
DELETE FROM Student  
WHERE ...;
```

Foreign keys + DELETE: What's the issue?

Student (id, first_name, last_name, city, state)

Enrollment (student_id, crn)

Course (crn, dept_code, number, title)

```
DELETE FROM Student  
WHERE id = 'S123456789';
```

Modifying & deleting data affects other tables

Student (id, first_name, last_name, city, state)

Enrollment (student_id, crn)

Course (crn, dept_code, number, title)

Need to maintain
referential integrity!

```
DELETE FROM Student  
WHERE id = 'S123456789';
```

By default, UPDATE &
DELETE fail when they would
break referential integrity.

Modifying & deleting data affects other tables

Student (id, first_name, last_name, city, state)

Enrollment (student_id, crn)

Course (crn, dept_code, number, title)

```
CREATE TABLE Enrollment (  
    ...  
    FOREIGN KEY crn REFERENCES Course (crn)  
    ON DELETE ...,  
    FOREIGN KEY student_id REFERENCES Student (id)  
    ON UPDATE ...  
);
```

- **RESTRICT**
- **CASCADE**
- NO ACTION
- SET NULL
- SET DEFAULT

Modifying data & natural vs. synthetic keys

Natural	<p>You can get a new SSN in cases of stolen identity.</p> <p>Account numbers can change when banks merge.</p> <p>Fits with ON UPDATE CASCADE.</p>
Synthetic	<p>Should never change.</p> <p>Fits with ON UPDATE RESTRICT / NO ACTION.</p>

Usually don't want to DELETE!

Soft delete

- I might want this data later, after all.
- Delete = Flag this data as “inactive”.
- E.g., closed customer account or discontinued product.
- Fits with ON DELETE RESTRICT / NO ACTION.


Hard delete

- I want to delete the data, so dang it, delete it!
- Fits with ON DELETE CASCADE.

Getting data from/to a file

Not a trivial thing!

- File on client or server?
- Security issues



Adding records to a table from a file

PostgreSQL

Server file: `COPY Product FROM '/my/path/product.csv' CSV HEADER;`


Must be logged in as a superuser, e.g., postgres. Server app must have file permission.

Client file: `\copy Product FROM '/my/path/product.csv' CSV HEADER`

Must use psql.

<u>p_name</u>	price	manufacturer
---------------	-------	--------------

p_name, price, manufacturer
Gizmo,19.99,GizmoWorks
Powergizmo,39.99,GizmoWorks
Widget,19.99,WidgetsRUs
HyperWidget,203.99,Hyper



<u>p_name</u>	price	manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	39.99	GizmoWorks
Widget	19.99	WidgetsRUs
HyperWidget	203.99	Hyper

Add, modify, delete from schema

Database schemas not static

Needs change over time.

Think about long-term plans & maintenance.



Add, modify, delete

Tables

```
CREATE TABLE Student (...);
```

```
ALTER TABLE Student  
RENAME TO Scholar;
```

```
DROP TABLE Student;
```

Attributes

```
ALTER TABLE Student  
ADD COLUMN zip INT;
```

```
ALTER TABLE Student  
RENAME COLUMN zip TO zip_code;
```

```
ALTER TABLE Student  
ALTER COLUMN zip VARCHAR(9);
```

```
ALTER TABLE Student  
DROP COLUMN state;
```


Constraints

Similar. Later.

Which can change existing query semantics?


- A. Adding a table
- B. Renaming a table
- C. Deleting a table
- ✓ D. All of the above

Should be obvious.



Which can change existing query semantics?

- A. Adding an attribute
- B. Renaming an attribute
- C. Deleting an attribute
- D. B&C
- ✓ E. All of the above



```
SELECT *  
FROM ...  
...;
```

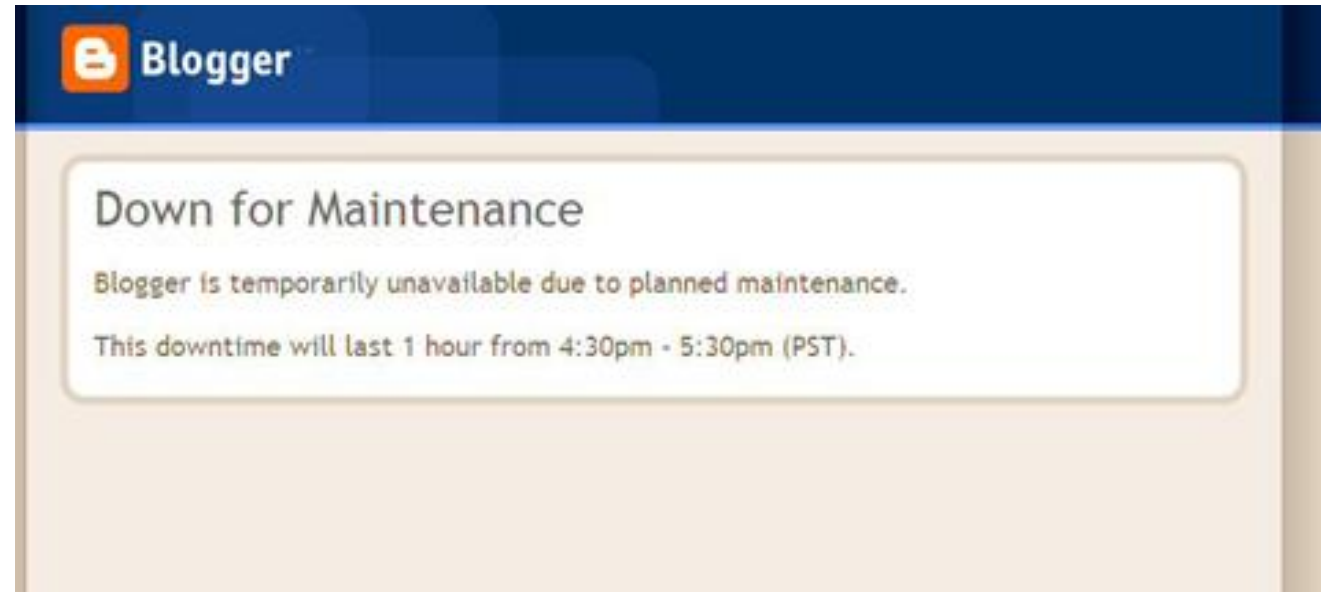
Dangerous – avoid SELECT *!
Specify columns explicitly.

Process for changing schemas

Traditional – Regular shutdowns

1. Make & test changes on a test server.
2. Have regularly scheduled system shutdown.
 - a) Typically late at night to minimize user disruption.
 - b) Disable user access to live server.
 - c) Make & test changes on the live server.

Takes time. Lots of data, possibly replicated & distributed.



Reducing downtime

Instead of a full shutdown, just block users from DB temporarily.

- Shutdown only as long as needed.
- “Lock” the DB.

Only need to lock tables that are changing.

- See COMP 322 & 421 for more about locks.

For UX, assumes lock is only needed briefly.

Updating without any user disruption

E.g.:

- Create “shadow table” incorporating changes
- Create “triggers” in original table that forward data updates to shadow table
- Copy original table’s data to shadow table
- Rename shadow table to replace original table

Can choose to only let some users see the changes.

Some DBMS provide support for such “online” schema changes.

Structured Query Language (SQL) Overview

Syntax:	Mostly declarative Old-fashioned, COBOL-based Numerous “standards”: ANSI SQL, SQL92, SQL99 Plus vendor-specific extensions & subsets
Semantics:	Mostly straight-forward Occasionally confusing
Implementation:	Highly optimized, parallelized Occasionally needs hints

