

COMP 430

Intro. to Database Systems

SQL from application code

Some issues

- How to connect to database
 - Where, what type, user credentials, ...
- How to send SQL commands
- How to get communicate data to/from DB
 - Object-Relational Impedance Mismatch
- What code goes where

Connecting to DB

Connecting to database – Java example

```
public class Example
{
    public static void main(String[] args)
    {
        Connection connection = null;
        Class.forName("com.Microsoft.jdbc.sqlserver.SQLServerDriver");
        String url = "jdbc:microsoft:sqlserver://localhost:1433;DatabaseName=MYDB";
        connection = DriverManager.getConnection(url, "USERNAME", "PASSWORD");
        ...
    }
}
```

Imports and error handling
omitted for brevity.

Key pieces:

- Driver
- Host
- DB name
- User credentials

Need sqljdbc4.jar in CLASSPATH.

Connecting to database – Python + SQLAlchemy example

String split for readability.

```
engine = create_engine("mssql+pyodbc://USERNAME:PASSWORD@localhost/MYDB" +  
                        "driver=SQL+Server+Native+Client+10.0")
```

Key pieces:

- Driver
- Host
- DB name
- User credentials


Connecting to database

Those *connection strings* share a standard syntax, although some arguments can be specified separately by library.

DB/Application Interface Overview



Mapping data is **HARD**



Mapping data is **HARD**

Under the hood, often involves two mappings:

- application → string
- string → DB

Communicating & (un)parsing data is relatively expensive.

Strategy: Plain SQL

```
connection = ...;  
Statement stmt1 = connection.createStatement();  
stmt1.executeUpdate("CREATE TABLE Student" +  
                    "(id INT, first VARCHAR(50), last VARCHAR(50))");  
...  
Statement stmt2 = connection.createStatement();  
ResultSet result = stmt2.executeQuery("SELECT id FROM Student");  
...
```

Plain SQL Summary

Minimal API

- ☺ SQL commands not restricted by API
- ☹ Requires knowing SQL
- ☹ Limited static type checking possible
- ☹ Dynamic SQL strings allow injection attacks



What are some bad input strings?

```
studentId = getQueryString("StudentID");  
String query = "SELECT first, last FROM Student WHERE id = " + studentId;  
Statement stmt = connection.createStatement(query);
```

Satisfy WHERE clause:

"123456789 OR 1=1"

Cause destructive behavior:

"123456789; DROP TABLE Student"

Expose implementation details:

"1 AND userid IS NULL"

Many variations on these themes.

Techniques for preventing injection attacks

- Validate input used in SQL command strings
 - E.g., numeric input only contains digits
- Build SQL commands via parameterized APIs (example)
- Only call DB functions; no permission to access tables directly
- Tightly manage user permissions
- Don't expose DB error messages

```
studentId = getRequestString("StudentID");
Statement stmt = connection.prepareStatement(
    "SELECT first, last FROM Student WHERE id=?");
Stmt.setInt(1, Integer.parseInt(studentId));
ResultSet result = Stmt.executeQuery(stmt);
while (result.next()) {
    String first = result.getString("first");
    String last = result.getString("last");
    ...
}
```

Strategy: Object-Relational Mapping Library

- Think in terms of application & persistent objects
 - Specify which objects persistent
 - Interaction with DB largely(?) hidden
- Basic abstraction: table = class, table row = ob
- Generates both from specification
 - Generates CRUD code
 - Maps between application/DB types
- Manages data movement



ORM: Table/object specification

Hibernate

```
<hibernate-mapping>
  <class name="Student" table="Student">
    <id name="id" type="int" column="id">
      <generator class="native"/>
    </id>
    <property name="firstName" column="first_name" type="string" />
    <property name="lastName" column="last_name" type="string" />
  </class>
</hibernate-mapping>
```

SQLAlchemy

```
class Student(Base):
    __tablename__ = 'Student'
    id = Column(Integer, primary_key=True, autoincrement=True)
    first_name = Column(String)
    last_name = Column(String)
```

ORM: Data movement

- Granularity:
 - Application-level traditionally accesses an attribute at a time.
 - DB-level accesses one or more records at a time.
- Consistency
 - Are application & DB data guaranteed to be consistent?
- Eager vs. lazy loading

ORM Summary: Pros

☺ Simplifies application programming

- Avoids SQL syntax
- Fits in style of language
- Allows static type checking

```
resultset = session.query(User, Document, DocumentPermission)
                .join(Document)
                .join(DocumentPermission)
                .filter(User.email == 'john_smith@foo.com')
```

ORM Summary: Cons

☹️ Doesn't really simplify learning curve

- Often largely SQL in disguise
- Still need to understand SQL concepts
- Complex library
- Every ORM is different.

☹️ Efficiency overhead

☹️ May encourage bad usage

- SELECTing all fields
- n+1 SELECTs problem

“ORMs make the easy even easier, and the difficult impossible. ... writing good, high-performing SQL is hard enough as it is. ... Add an ORM ... and then you have to figure out how to say what you want ... through a mapping layer.”

n+1 SELECTs Problem

ORM pseudo-code:

```
foreach emp in select_all_employees()  
    print select_salary_of_employee(emp)
```

Avoid looping
over SELECT!

Corresponds to:

```
SELECT * FROM Employee;
```

```
SELECT salary FROM Salary WHERE emp_id = 1
```

```
SELECT salary FROM Salary WHERE emp_id = 2
```

```
...
```

Instead of:

```
SELECT salary  
FROM Salary s
```


```
INNER JOIN Employee e ON e.emp_id = s.emp_id
```

Strategy: Write your own ORM

😊 Specialized to what you need

😞 Lots of code

😞 Error-prone



Where does “business logic” go?



What's this “business logic”?

Common term for the core computation.

Example for a bank:

Business logic:

- Creating accounts
- Computing interest
- Transactions between accounts

Application logic:

- Website/mobile UI
- User log-in

Logic in application only

DB as dumb storage. Even data-integrity rules in application.

- Add application servers to handle load.
- Support multiple DBMS's (at once or over time)
 - E.g., creating application s/w that is DB-agnostic
 - DBMS's are commodities. Use whichever is currently best.
 - Put logic in the single application, rather than replicating in multiple DBMS's.
- SQL language & tools (e.g., unit testing) aren't as good.
- Decision-makers & software engineers are application-focused.
 - Few good SQL programmers.
 - DB programmers often in a separate DB team.

Logic in DB only

DB in charge of data – including data integrity, data computation

- Use functions, triggers extensively
 - Application accesses DB via DB-level function API
 - Add DB servers & replication to handle load
-
- Application only gets the data it really needs
 - Minimize communication.
 - Expose minimum of DB data & structure. Maximize security.
 - Support multiple applications (at once or over time) using same DB
 - Applications change more frequently than DB architectures.
 - Application programmers shouldn't write SQL or pseudo-SQL (ORM)
 - Protect data from bad programmers.
 - Maximize query optimization.

Logic trade-off summary

Application:

- **Convenience** – More Java/C++/... programmers than for SQL. Better language, libraries, & tools.
- **Flexibility** – Can use the DB however you want.

DB:

- **Encapsulation** – Application programmer only uses API, and doesn't need details, including much SQL.
- **Safety** – Protect DB from application programmer. Only expose desired API.
- **Performance** – Less I/O. Better DBMS optimization.

Between the extremes

Application & DB each used for their best features

- DB – CRUD, data integrity, data replication
- Business logic in either application or DB, **depending on situation.**
- Application – Everything else

<http://www.vertabelo.com/blog/notes-from-the-lab/business-logic-in-the-database-yes-or-no-it-depends>