

COMP 430/533

Intro. to Database Systems

Indexing

How does DB find records quickly?

- Various forms of *indexing*
- An index is automatically created for primary key.
- SQL gives us some control, so we should understand the options.
 - Mainly concerned with user-visible effects, not underlying implementation.

```
CREATE INDEX index_city_salary  
ON Employees (city, salary);
```

```
CREATE UNIQUE CLUSTERED INDEX idx  
ON MyTable (attr1 DESC, attr2 ASC);
```

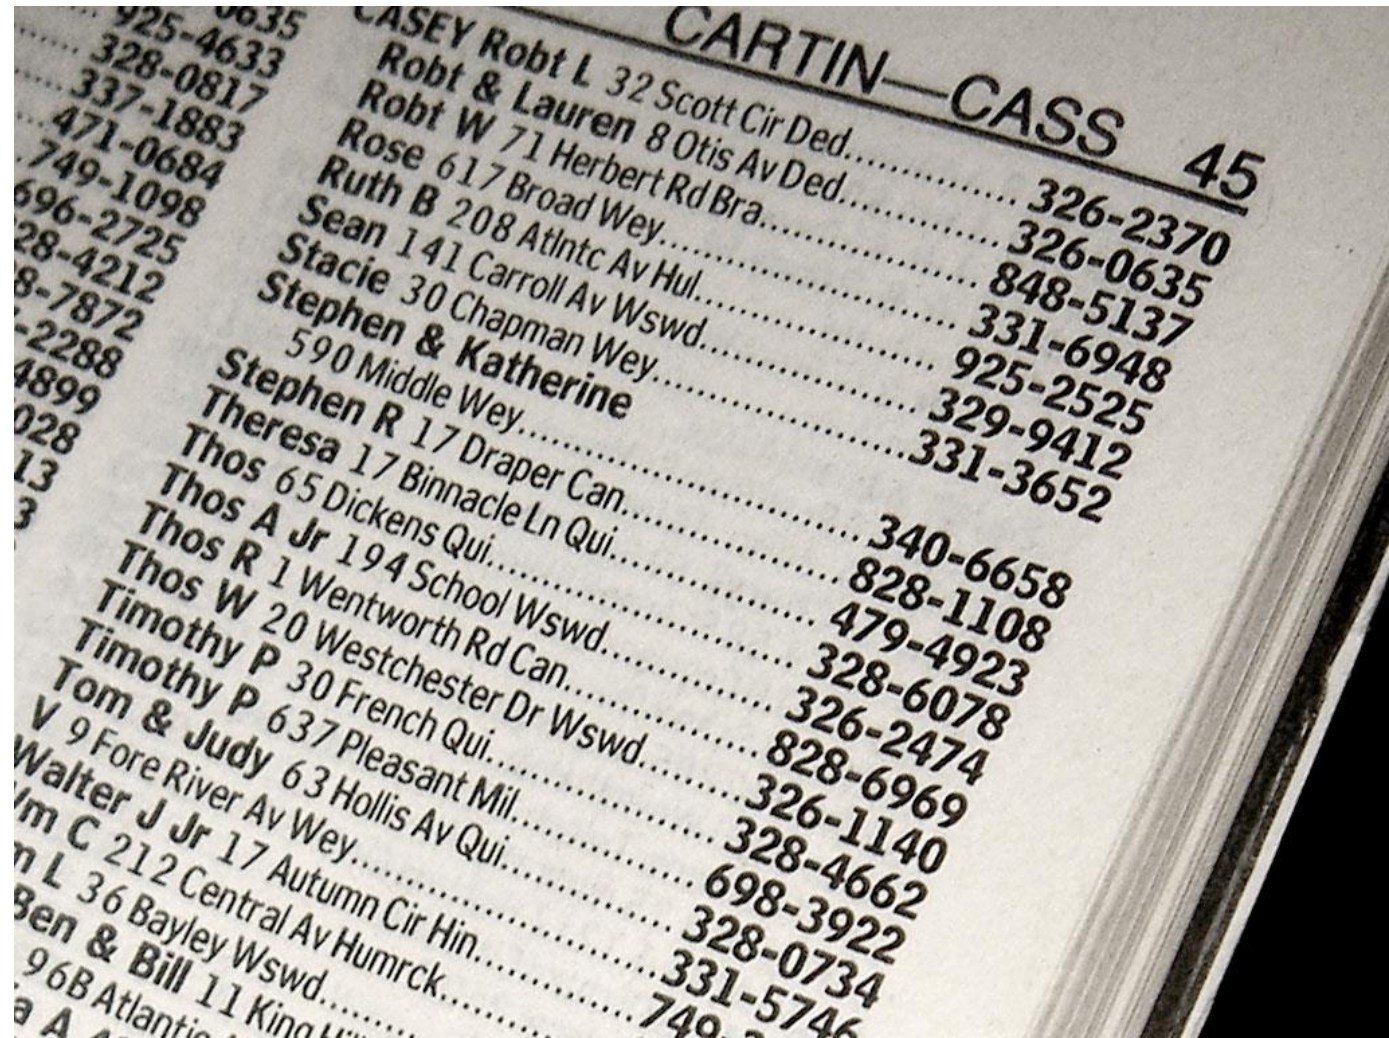
Options vary.
Not PostgreSQL.

Phone book model

Data is stored with search key.
Clustered index.

Organized by **one search key**:

- Last name, first name.
- Searching by any other key is slow.




Library card catalog model

Index stores pointers to data.

Non-clustered index.

Organized by search key(s).

- Author last name, first name.
- Title
- Subject



Evaluating an indexing scheme


- Access type flexibility
 - Specific key value – e.g., “John”, “Smith”
 - Key value range – e.g., salary between \$50K and \$60K

Advantage:

- Access time

Disadvantages:

- Update time
- Space overhead



Creating an index can slow down the system!

Ordered indices

Sorted, as in previous human-oriented examples

Types of indices we'll see

- Clustered vs. non-clustered
- Dense vs. sparse
- Unique data vs. not-unique
- Single- vs. multi-level

These ideas can be combined in various ways.

Clustered index – dense or sparse

No index

Data		
	10	
	10	
	20	
	30	
	30	
	100	

...

Dense index

Index		Data		
10		→		10
20		→		10
30		→		20
		→		30
100		→		30
		→		100

...

...

Sparse index

Index		Data		
10		→		10
30		→		10
		→		20
		→		30
		→		30
		→		100

...

...

- Common for primary keys
- Not supported in PostgreSQL
- Sparse – typically pointer to top of each file block

- Faster access?
- Faster insert, update, delete?
- Less space?

CARTIN—CASS 45


325-4613	RSEY Robt L 32 Scott Cir Ded	326-2370
326-0817	Robt & Lauren 8 Ovis Av Ded	326-0635
337-1883	Robt W 71 Herbert Rd Bra	848-5137
471-0804	Ruth B 208 Alente Wey	331-6948
749-1098	Sean 141 Carroll Av Hul	925-2525
680-2725	Stacie 30 Chapman Wey	329-9412
28-4212	Stephen & Katherine	331-3652
8-7872	Stephen R 17 Draper Can	340-6658
2288	Theresa 17 Binnacle Ln Qui	828-1108
899	Thos 65 Dickens Qui	479-4923
028	Thos A Jr 194 School Wswd	328-6078
13	Thos W 20 Westworth Rd Can	326-2474
3	Timothy P 30 French Qui	328-6969
	Tom & Judy 637 Pleasant Mil	326-1140
	V 9 Fore River Av Wey	328-4662
	Walter J Jr 17 Autumn Cir Hin	698-3922
	m C 212 Central Av Huntrck	328-0734
	m L 30 Bayley Wswd	331-5746
	Ben & Bill 11 Kim	740
	908 Atlant	
	9 A Atlant	

PostgreSQL & clustering

- Clustered index not supported
- Can cluster data (i.e., physically sort) for efficient access.
 - This sorted order is **not** maintained after CRUD operations.

```
CLUSTER Employees USING index_city_salary;
```

Simple non-clustered index – must be dense





So far, index size proportional to #values

For efficient access, want index to fit in memory.

Solution: Multi-level index

Same problem & solution as for page tables in virtual memory.

Multi-level indices



CRUD on dense indices or clustered data

	10	
	20	
	30	
	40	
	50	
	60	
	70	

...



	10	
	20	
	25	
	30	
	40	
	50	
	60	
	70	

...

Moving all records is too expensive.
Otherwise, results in fragmentation.

Fragmentation consequences

Fragmentation (empty spaces) increases with usage.

Performance degrades with usage.

Need to periodically reorganize data & indices.

Solution: B+-trees

Ordered indices – B+-trees

The most commonly used indexing algorithm


B+-trees

- Balanced search trees optimized for disk-based access
 - Reorganizes a little on every update
 - Shallow & wide (high fan-out)
 - Typically, node size = disk block
 - Slight differences from more commonly-known B-trees
 - All data in leaf nodes
 - Leaf nodes sequentially linked
- } Easily get all data in order.

```
CREATE INDEX ... ON ... USING BTREE;
```

Or, is the default
in many DBMSs.

B+-tree example



... Data records ...

B+-tree performance

- Very similar to multi-level index structure
 - Slightly higher per-operation access & update time
 - + No degradation over time
 - + No periodic reorganization

B+-tree widely used in relational DBMSs

What about non-unique search keys?

- Can allow duplicates in tree.
 - Maintain \leq order instead of $<$.
 - Complicates details.
- Can either return first or all duplicates.
 - Easy, since duplicates are adjacent.


Indexing on VARCHAR keys

- Key size variable, so number of keys that fit into a node also varies.
- Techniques to maximize fan-out:
 - Key values at internal nodes can be prefixes of full key. E.g., “Johnson” and “Jones” can be separated by “Jon”.
 - Key values at leaf nodes can be compressed by sharing common prefixes. E.g., “Johnson” and “Jones” can be stored as “Jo” + “hnson”/”nes”

Hash indices

An alternative to ordered indices

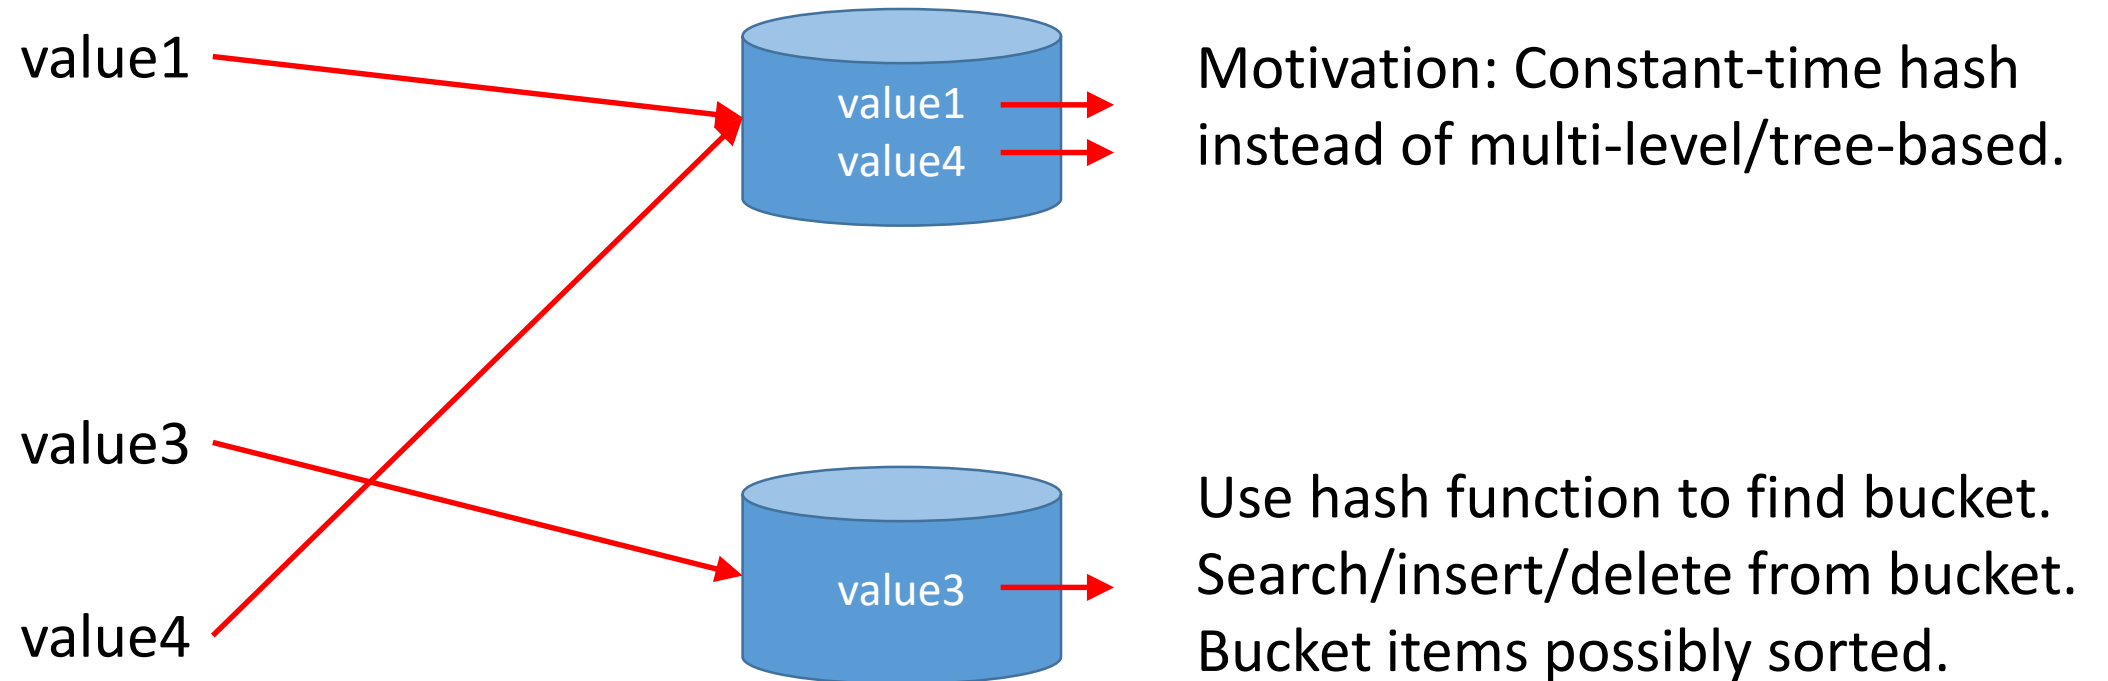
Basic idea of a hash function



$h()$ distributes values uniformly among the buckets.

$h()$ distributes typical subsets of values uniformly among the buckets.

Using hash function for indexing



Buckets can overflow

- Overflow some buckets – skewed usage
- Overflow many buckets – not enough space reserved

Solutions:

- Chain additional buckets – degrades to linear search
- Stop and reorganize
- Dynamic hashing (extensible or linear hashing) – techniques that allow the number of buckets to grow without rehashing existing data

Advantages & disadvantages of hash indexing

- + One hash function vs. multiple levels of indexing or B+-tree
- Storage about the same.
 - + Don't need multiple levels.
 - But, need buckets about half empty for good performance.
- Can't easily access a data range. Thus, poor for relations like $>$.
- Simple hashing degrades poorly when data is skewed relative to the hash function.

Multiple indices & Multiple keys

Using indices for multiple attributes

What if queries like the following are common?

```
SELECT ...  
FROM Employee  
WHERE job_code = 2 AND performance_rating = 5;
```

Strategy 1 – Index one attribute

```
CREATE INDEX idx_job_code on Employee (job_code);  
  
SELECT ...  
FROM Employee  
WHERE job_code = 2 AND performance_rating = 5;
```

Internal strategy:

1. Use index to find Employees with job_code = 2.
2. Linear search of those to check performance_rating = 5.

Strategy 2 – Index both attributes

```
CREATE INDEX idx_job_code on Employee (job_code);  
CREATE INDEX idx_perf_rating on Employee (performance_rating);  
  
SELECT ...  
FROM Employee  
WHERE job_code = 2 AND performance_rating = 5;
```

Internal strategy – chooses between

- Use job_code index, then linear search on performance_rating.
- Use performance_rating index, then linear search on job_code.
- Use both indices, then intersect resulting sets of pointers.

Strategy 3 – Index attribute set

```
CREATE INDEX idx_job_perf on Employee (job_code, performance_rating);  
  
SELECT ...  
FROM Employee  
WHERE job_code = 2 AND performance_rating = 5;
```

Attribute sets ordered lexicographically:

$(jc1, pr1) < (jc2, pr2)$ iff either

- $jc1 < jc2$
- $jc1 = jc2$ and $pr1 < pr2$

This strategy typically
uses ordered index,
not hashing.

Note that this prioritizes job_code over performance_rating!

Strategy 3 – Index attribute set

Efficient


```
CREATE INDEX idx_job_perf on Employee (job_code, performance_rating);  
  
SELECT ... FROM Employee WHERE job_code = 2 AND performance_rating = 5;  
  
SELECT ... FROM Employee WHERE job_code = 2 AND performance_rating < 5;
```

Inefficient

```
CREATE INDEX idx_job_perf on Employee (job_code, performance_rating);  
  
SELECT ... FROM Employee WHERE job_code < 2 AND performance_rating = 5;  
  
SELECT ... FROM Employee WHERE job_code = 2 OR performance_rating = 5;
```

Strategy 4 – Grid indexing

- n attributes viewed as being in n -dimensional grid
- Numerous implementations
 - Grid file, K-D tree, R-tree, ...
- Mainly used for spatial data
- GIS DBs are a separate, specialized topic.
 - Many RDBMS's have a GIS-specialized component.



Strategy 5 – Bitmap indices

- Coming next...

Bitmap indices

One more alternative approach for indexing

Basic idea of bitmap indices

Assume records numbered 0, 1, ..., and easy to find record $#i$.

Record #	id	gender	income_level
0	51351	M	1
1	73864	F	2
2	13428	F	1
3	53718	M	4
4	83923	F	3

Bitmaps

M	F	IL1	IL2	IL3	IL4	IL5
1	0	1	0	0	0	0
0	1	0	1	0	0	0
0	1	1	0	0	0	0
1	0	0	0	0	1	0
0	1	0	0	1	0	0

F's bitmap: 01101

Queries use standard bitmap operations

```
SELECT ... FROM ... WHERE gender = 'F' AND income_level = 1;
```

F's bitmap		1's bitmap	
01101	&	10100	= 00100

```
SELECT ... FROM ... WHERE gender = 'F' OR income_level = 1;
```

F's bitmap		1's bitmap	
01101		10100	= 11101

```
SELECT ... FROM ... WHERE income_level <> 1;
```

~ 1's bitmap	
10100	= 01011

Space overhead

- Normal: $\#records \times (\#values + 1)$ bits, if nullable
 - Typically used when $\#$ of values for attribute is low.
- Encoded: $\#records \times \log(\#values)$
 - But need to use more bitmaps
- Compressed: $\sim 50\%$ of normal
 - Can do bitmap operations on compressed form.

A couple details

- When deleting records, either ...
 - Delete bit from every bitmap for that table, or
 - Have an *existence bitmap* indicating whether that row exists.
- Can combine bitmaps with B+-trees.
 - B+-tree leaf holds list of pointers to records with a particular attribute value.
 - Can instead hold bitmap for that value.

Summary of index types

Single-level	Index generally too large to fit in memory
Multi-level	Degrades due to fragmentation
B+-tree	General-purpose choice
Hash	Good for equality tests; potentially degrades due to skew & overflow
Spatial	Good for GIS/spatial coordinates
Bitmap	Good for multiple attributes each with few values