

# COMP 430

# Intro. to Database Systems

Query optimization

# Questions we want to address

- How to understand queries' performance?
- What queries can system optimize for us?
  - How can we tell?
  - ~~How does it optimize?~~
- What queries do we need to optimize ourselves?
  - How can we tell?
  - How do we optimize?

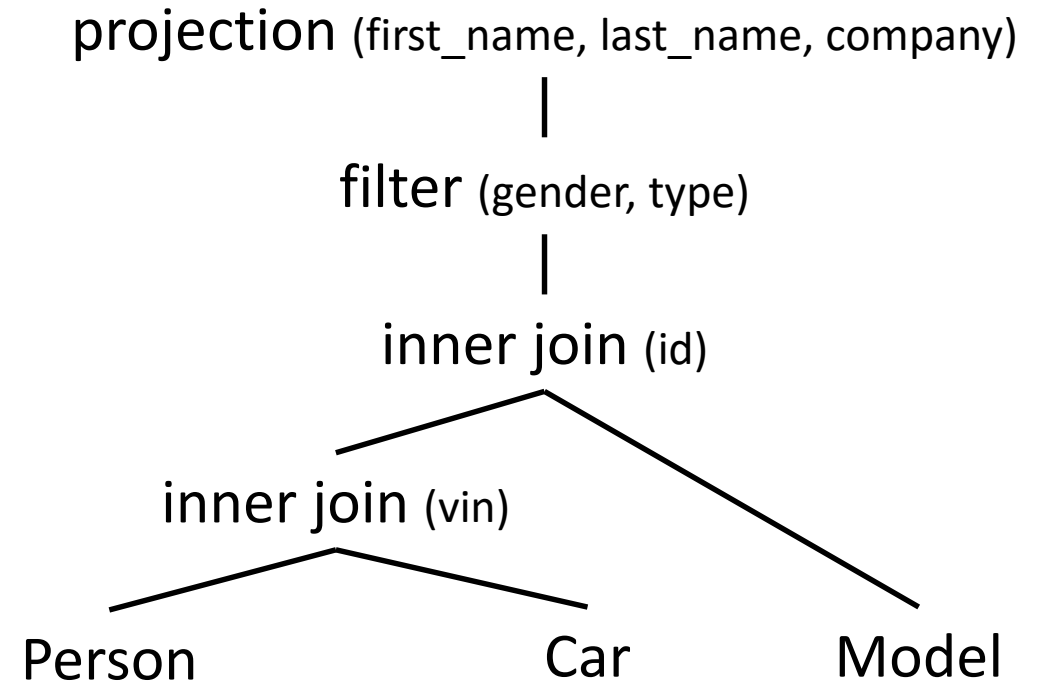
Need to understand some basic optimization strategies, but not the details.

# Example query

```
SELECT Person.first_name, Person.last_name, Model.company  
FROM Person  
INNER JOIN Car ON Person.car_vin = Car.vin  
INNER JOIN Company ON Car.model_id = Model.id  
WHERE Person.gender = 'F' AND Car.type IN ('sedan', 'coupe');
```

# A simple query tree

```
SELECT Person.first_name, Person.last_name, Model.company  
FROM Person  
INNER JOIN Car ON Person.car_vin = Car.vin  
INNER JOIN Company ON Car.model_id = Model.id  
WHERE Person.gender = 'F' AND Car.type IN ('sedan', 'coupe');
```

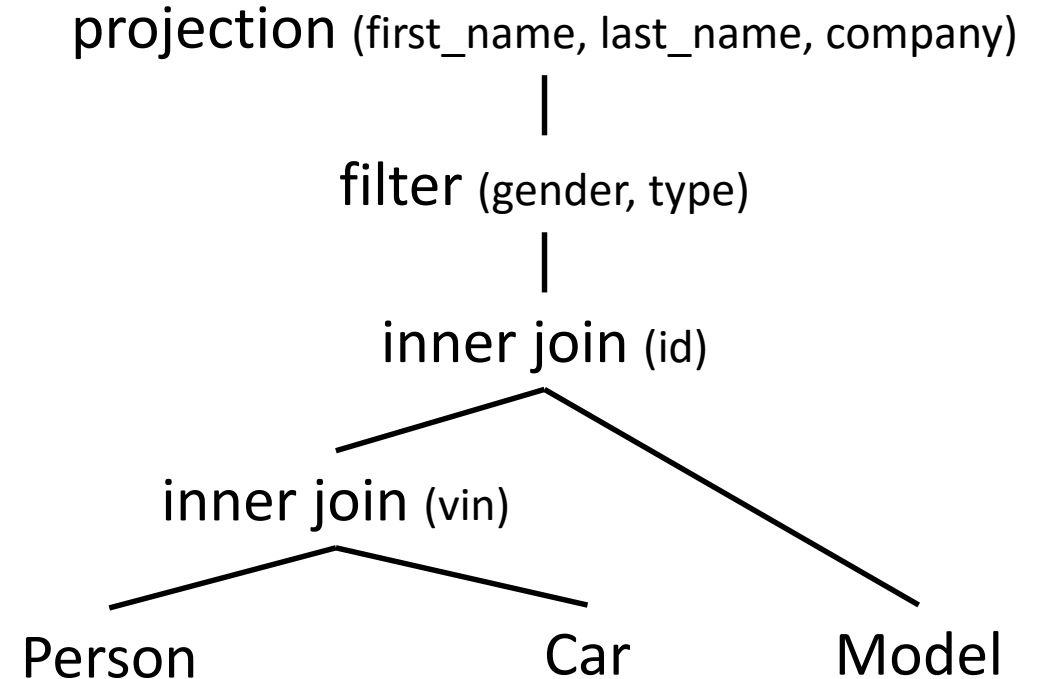


Tree form of *relational algebra* expression

# Activity – Generate equivalent query trees

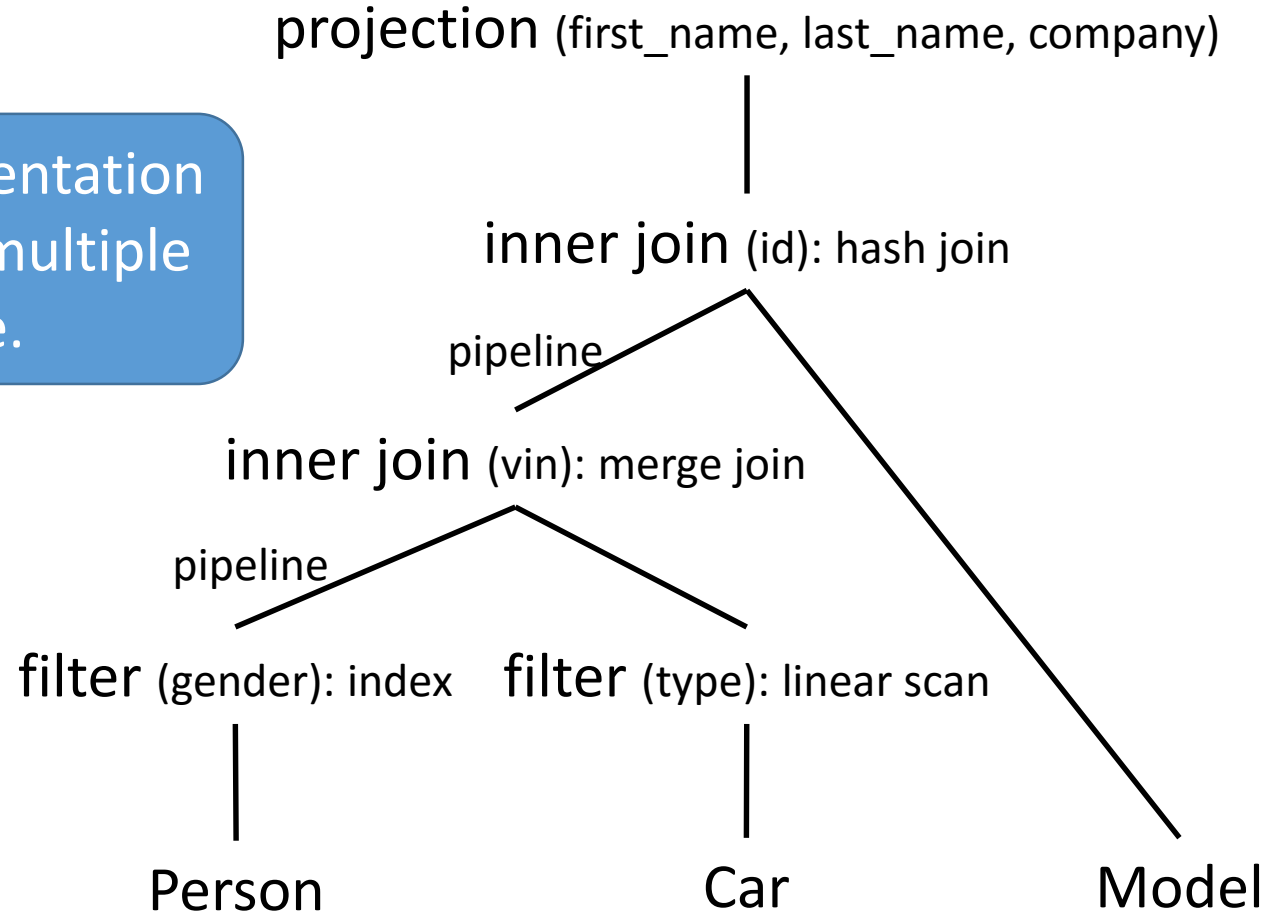
```
SELECT Person.first_name, Person.last_name, Model.company  
FROM Person  
INNER JOIN Car ON Person.car_vin = Car.vin  
INNER JOIN Company ON Car.model_id = Model.id  
WHERE Person.gender = 'F' AND Car.type IN ('sedan', 'coupe');
```

Suggestions?



# Generate execution plans

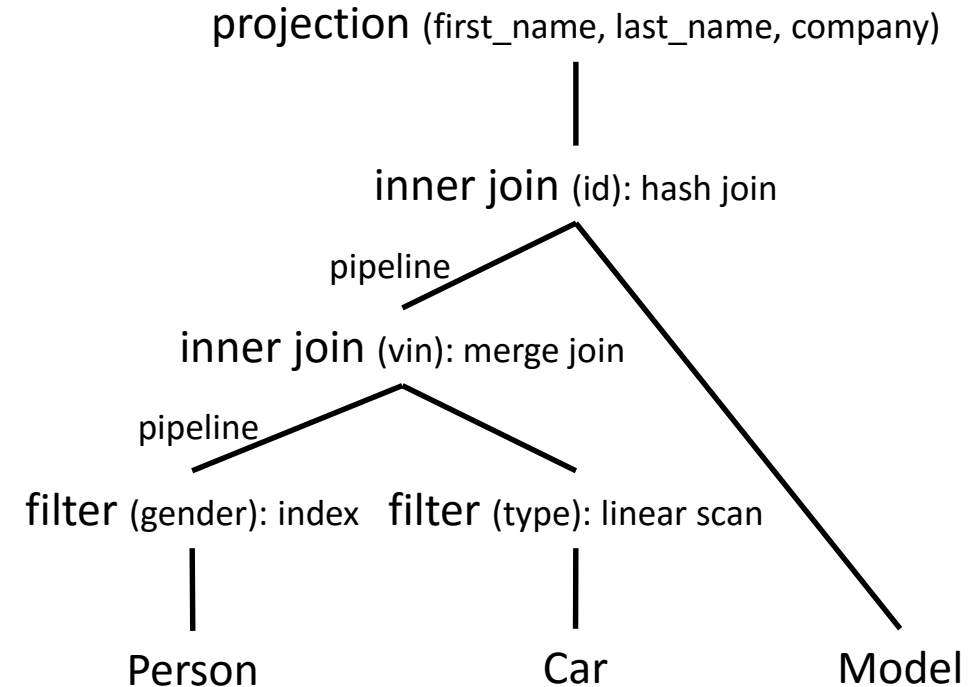
Notate with implementation methods. Possibly multiple plans per tree.



# Analyze costs of each plan

- Algorithm cost for each node
- Data amount
  - #tuples
  - Total #bytes or #disk blocks
  - Selectivity of filters
- Data access
  - Method – index, linear scan, binary search
  - Source – memory, disk

Data-dependent – determined at run time.  
Use & cache lowest cost plan.



# Trees → Plans → Costs

## Steps integrated

- Can be too many trees/plans to generate all of them.
- Cost heuristics guide tree/plan generation.



# Simple example of cost analysis

```
SELECT *  
FROM Person  
WHERE name IN ('Greiner', 'Halverhout') AND  
       zipcode = 77055 AND  
       birthdate > '19670212';
```

Test	CPU cost/record
name IN	100ns
zipcode =	1ns
birthdate >	1000ns

400,000,000 records

# Simple example of cost analysis: Test order

For each Person record:

If name IN ...

Then if zipcode = ...

Then if birthdate > ...

Then output record



How many possible orders?

How to calculate cost for each?

# Simple example of cost analysis: Test order

400,000,000 records

Test	CPU cost/record	Selectivity	Cost for this test	...simplified
name IN	100ns	0.00001	$400,000,000 \times 100\text{ns}$	40,000,000,000ns
zipcode =	1ns	0.0001	$400,000,000 \times 0.00001 \times 1\text{ns}$	4,000ns
birthdate >	1000ns	0.3	$400,000,000 \times 0.00001 \times 0.0001 \times 1000\text{ns}$	400ns
				<b>40,000,004,400ns</b>

Test	CPU cost/record	Selectivity	Cost for this test	...simplified
zipcode =	1ns	0.0001	$400,000,000 \times 1\text{ns}$	400,000,000ns
name IN	100ns	0.00001	$400,000,000 \times 0.0001 \times 100\text{ns}$	400,000ns
birthdate >	1000ns	0.3	$400,000,000 \times 0.00001 \times 0.0001 \times 1000\text{ns}$	400ns
				<b>400,400,400ns</b>

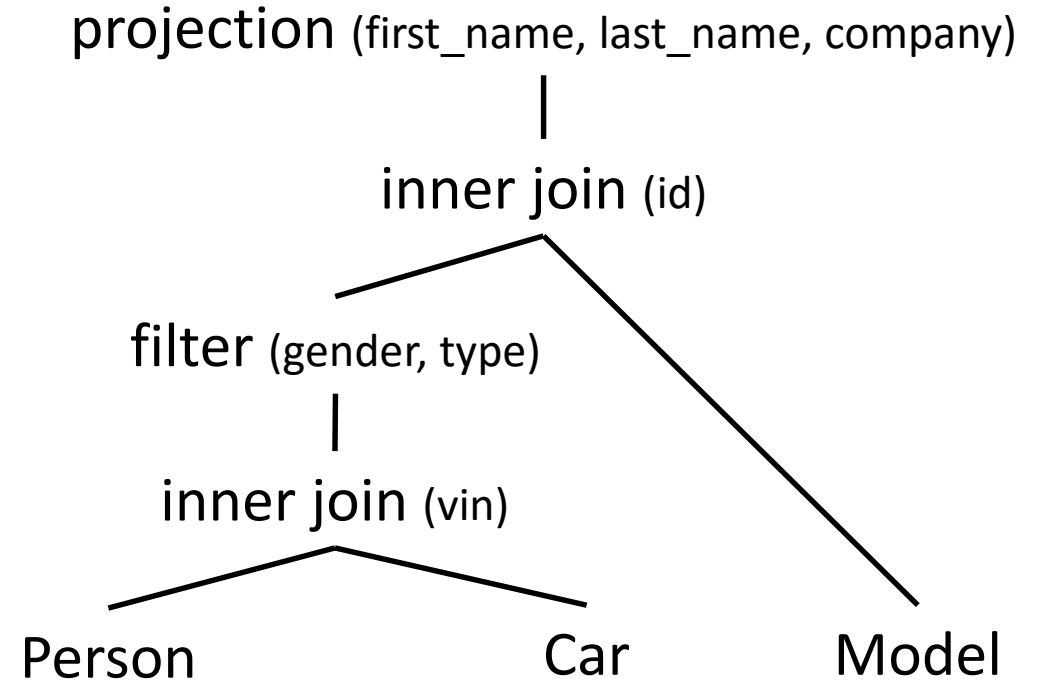
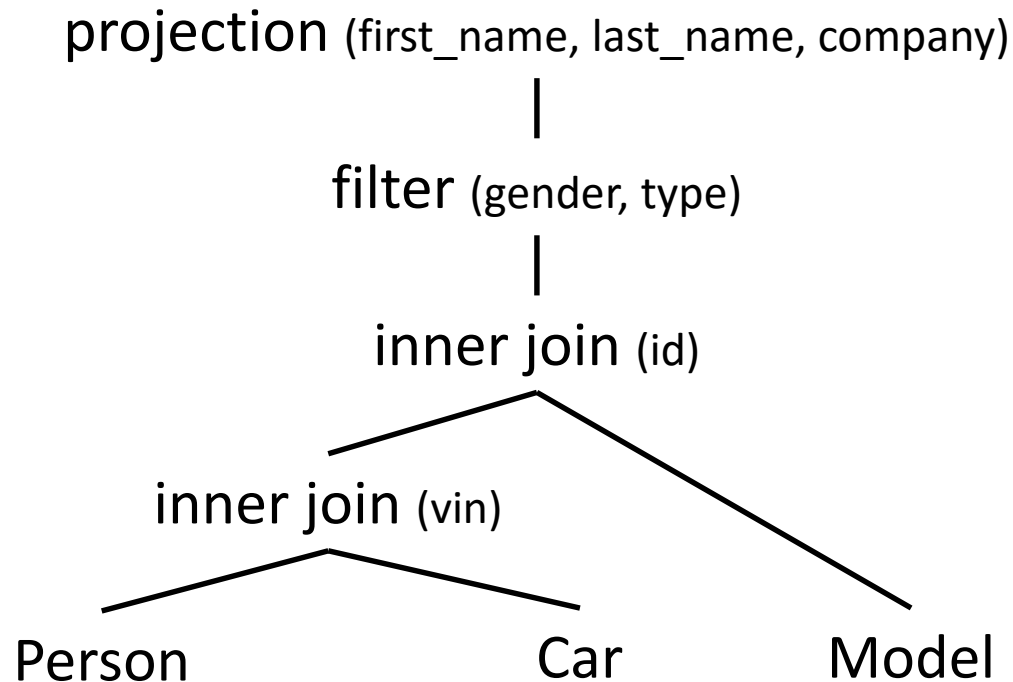
# Simple example of cost analysis: Test order

How to search through all the plans efficiently?

**Best:** in decreasing order by  $\frac{1 - \text{Selectivity}(\text{predicate})}{\text{Cost}(\text{predicate})}$ ,  
i.e., the fraction of data eliminated per unit time.

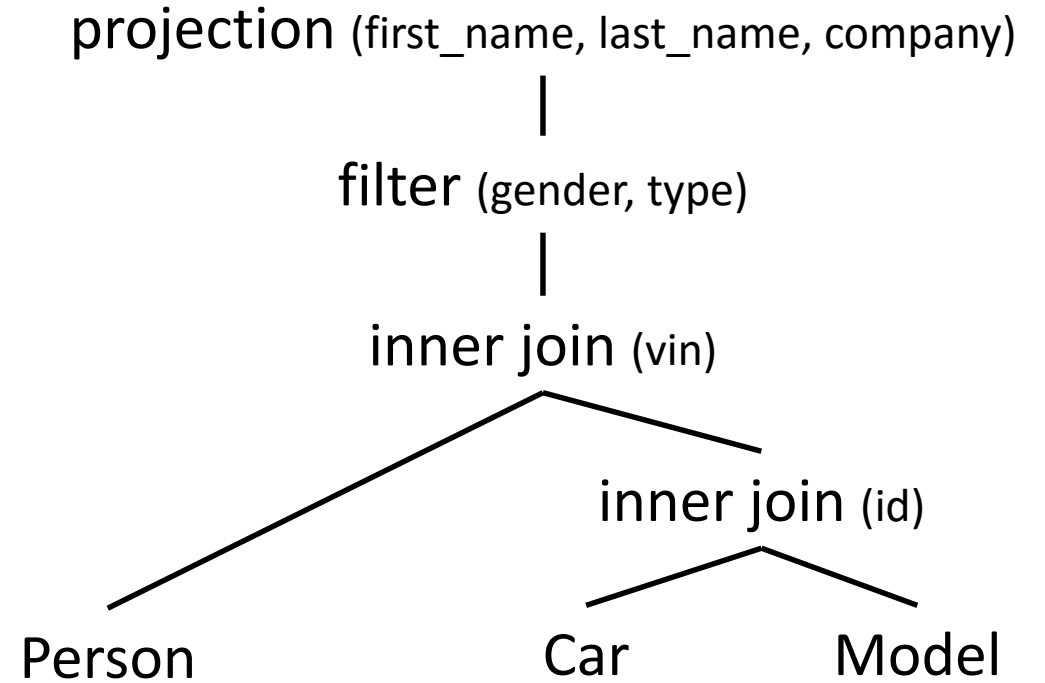
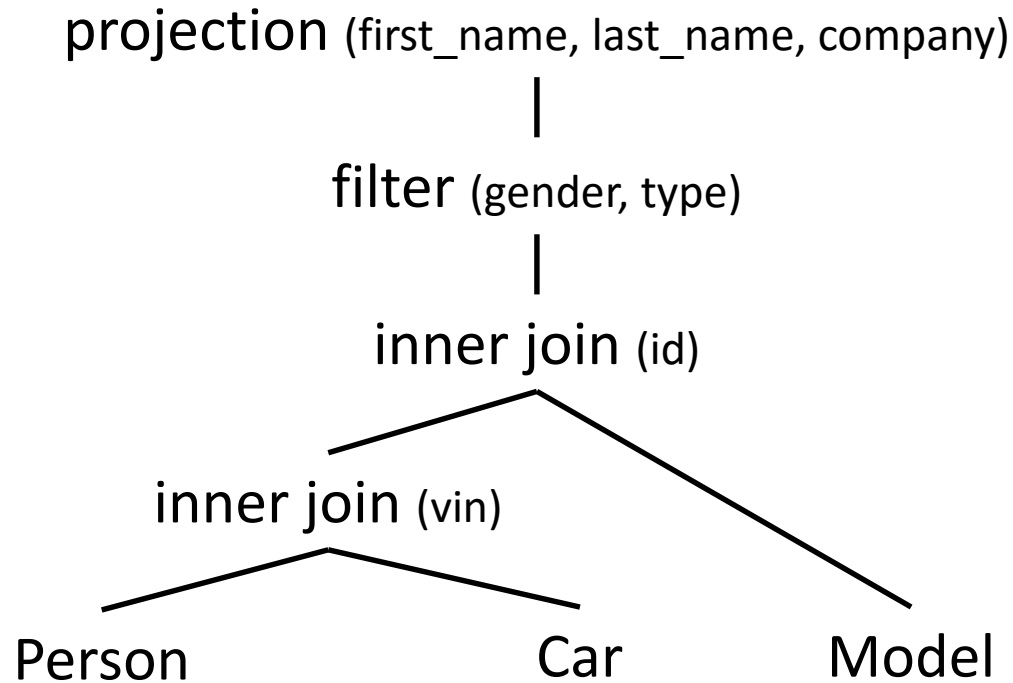
Transforming query trees

# Relational algebra equivalences: example



$$\text{filter}(\text{inner join}(A, B)) = \text{filter}_{AB}(\text{inner join}(\text{filter}_A(A), \text{filter}_B(B)))$$

# Relational algebra equivalences: example



Inner joins are associative & commutative.

# Too many trees to enumerate all

Example: inner join of  $n$  tables with commutativity & associativity  $\rightarrow$   
how many trees?

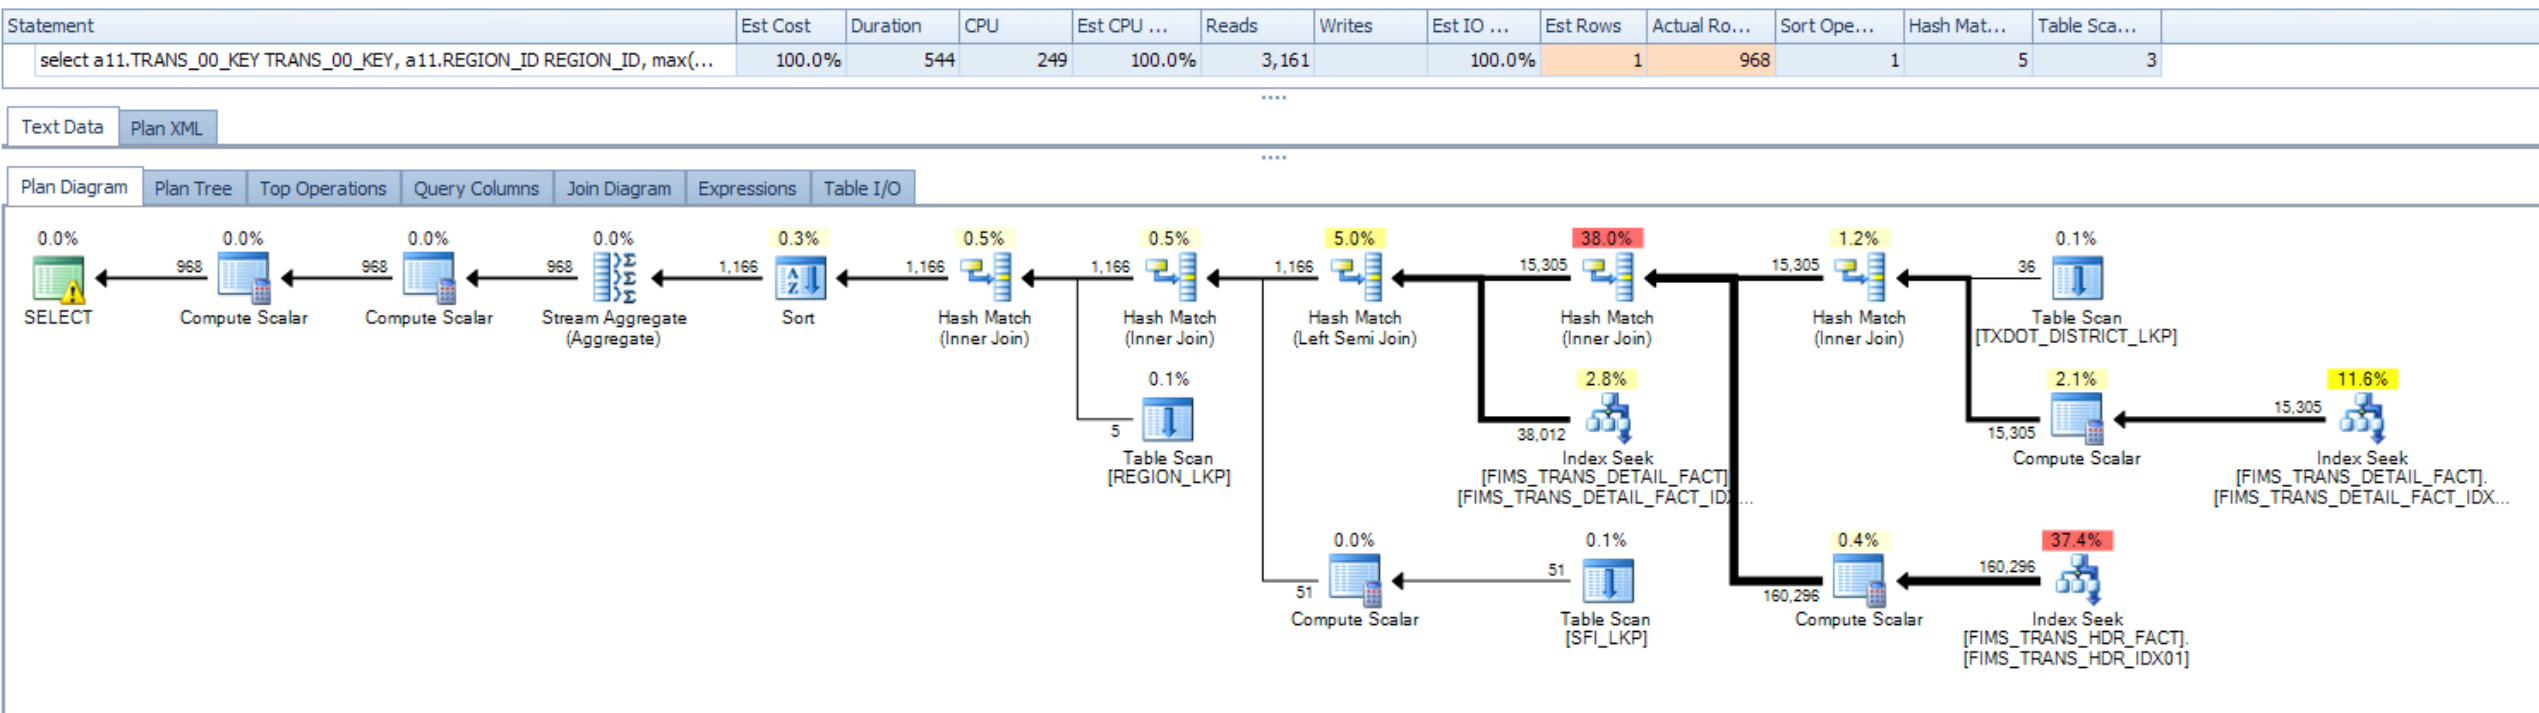
$$n! 2^n$$



# Implementation Strategies

# Viewing & understanding execution plan

Helpful to be able to understand plan and any bottlenecks.



# Scan implementations

- Indexing – previously discussed
- Binary search – if sorted on search key
- Linear search

# What to do with each operation's results?

- *Materialization*

- Data saved to disk (or memory, if small enough)

- Disk access

- + Can reuse data if needed multiple times

- **Pipelining**


- Consumer uses data as producer generates it

- Use buffer to allow consumer to be faster than producer

- + No disk access

# Left-deep trees have one long pipeline

Common to consider only *left-deep* trees of joins.



# Join implementations

Join(R, S) on  $R.A = S.A$

Let  $T(R) = \text{\#tuples in } R$ ,  $B(R) = \text{\#blocks in } R$

- Three main strategies
  - Nested joins + 2 variants
  - Sort-merge joins + 1 variant
  - Hash joins

# Nested loop join

```
For r in R:  
  For s in S:  
    If r.a == s.a, yield (r,s)
```

$$\text{Cost} = B(R) + T(R)B(S) + B(\text{result})$$

```
For s in S:  
  For r in R:  
    If r.a == s.a, yield (r,s)
```

$$\text{Cost} = B(S) + T(S)B(R) + B(\text{result})$$

# Block nested loop join – I/O aware

Let  $B+1 = \text{\#blocks in memory}$

```
For each group of  $B-1$  blocks  $br$  of  $R$ :  
  For block  $bs$  of  $S$ :  
    For  $r$  in  $br$ :  
      For  $s$  in  $bs$ :  
        If  $r.a == s.a$ , yield  $(r,s)$ 
```

$$\text{Cost} = B(R) + \frac{B(R)}{B-1} B(S) + B(\text{result})$$

Choose  $R$  = smaller relation



# Index nested loop join

Assumes S indexed on A.

For r in R:  
  If r.a == index(s,a), yield (r,s)

$$\text{Cost} = B(R) + T(R)C + B(\text{result})$$


$C$  = I/O to access all distinct values in index.  
Not constant, but typically  $< 10$ .

# Sort-merge join

1. Sort R and S each on A.
  - Use *external merge sort* (disk-based).
  - But, R and/or S might already be sorted on A.
2. “Merge” R and S.
  - Linear scan R.
  - Linear scan S, backing up on duplicates.

$$R \left( \log_B \frac{R}{B+1} \right) \\ S \left( \log_B \frac{S}{B+1} \right)$$


$$B(R) + B(S) \quad \dots \quad B(R)B(S)$$



Not too sensitive to skew.  
Leaves data sorted!

# Sort-merge join – I/O aware


Integrate the mergesort's merging with the “merge” of R,S




# Hash-join

Use hashing to decompose one large join into many small joins.

Step 1: Hash R, S each on same hash function into B buckets.



Step 2: Join pairs of buckets.  
Use hash-join recursively or  
block nested join.



Highly  
parallelizable

# Size estimation

Costs depend on amount of data

# Estimating sizes

- Size of each relation
- Selectivity of each join & filter condition
- How calculated?
  - Gather statistics
  - Use statistics to estimate selectivity

# Size statistics per relation $R$


- $tuples(R)$ : #tuples
- $length(R)$ : #bytes per tuple
- $bfactor(R)$ : #tuples that fit into a block (blocking factor)
- $blocks(R)$ : #blocks – if no fragmentation,  $= \left\lceil \frac{tuples(R)}{bfactor(R)} \right\rceil$

# Value statistics for attribute A

- $values(A)$ : #distinct values
- $minval(A)$ : minimum value
- $maxval(A)$ : maximum value
- Distribution of values for attribute A – estimated by histogram



# Equi-width histogram example



# Histograms summary

- Various other types of histograms
- High-level goals
  - Time- & Space-efficient to compute
  - Useful in estimating selectivity of equality and range conditions

# Managing statistics

- Accuracy vs. resources
  - Update frequency
  - Build histograms for which attributes, with how much accuracy
- Automatically by DBMS
- Some manual control
  - SQL Server: CREATE STATISTICS, UPDATE STATISTICS
- Can view statistics
  - SQL Server: DBCC SHOW\_STATISTICS

# Size estimates for various operations

- CrossJoin(R,S)      $tuples(R) \times tuples(S)$
  - InnerJoin(R,S)      $tuples(R)$
  - OuterJoin(R,S)      $tuples(R) + tuples(S)$
- } Assuming join on R's FK & S's PK
- 
- $R \cup S$       $tuples(R) + tuples(S)$
  - $R \cap S$       $\min(tuples(R), tuples(S))$
  - $R - S$       $tuples(R)$

# Calculating selectivity on filter conditions

- WHERE R.A = v:

- Without histogram:  $s = \frac{1}{values(R.A)}$

- With equal-width histogram:  $s = \frac{count(range\_containing\_A,R)}{tuples(R) \times range\_width}$

# Calculating selectivity on filter conditions

- WHERE R.A = S.B

- Without histograms:  $s = \min\left(\frac{1}{\text{values}(R.A)}, \frac{1}{\text{values}(S.B)}\right) = \frac{1}{\max(\text{values}(R.A), \text{values}(S.B))}$

- With histograms: Generalize the previous.

R.A values(R.A)=3	S.B values(S.B)=2	Match?
1	1	Yes
2	2	Yes
3	1	No
1	2	No
2	1	No
3	2	No

I.e., expect 1/3 to match.

# Calculating selectivity on filter conditions

- WHERE  $R.A \geq v$ :
  - Without statistics:  $s = \frac{1}{2}$
  - Without histogram:  $s = \begin{cases} 0 & \text{if } v > \text{maxval}(R.A) \\ \frac{\text{maxval}(R.A) - v}{\text{maxval}(R.A) - \text{minval}(R.A)} & \text{otherwise} \end{cases}$
  - With equal-width histogram: Generalize the previous.

# Calculating selectivity on filter conditions

- Conjunction:  $s = \prod s_i$
- Disjunction:  $s = 1 - \prod(1 - s_i)$
- Negation:  $s = 1 - s'$



# User-level Query Optimization

So, what should **you** the SQL programmer do?

# Just a bit of terminology

- Query optimization – what the system does
- Query tuning – what the programmer does

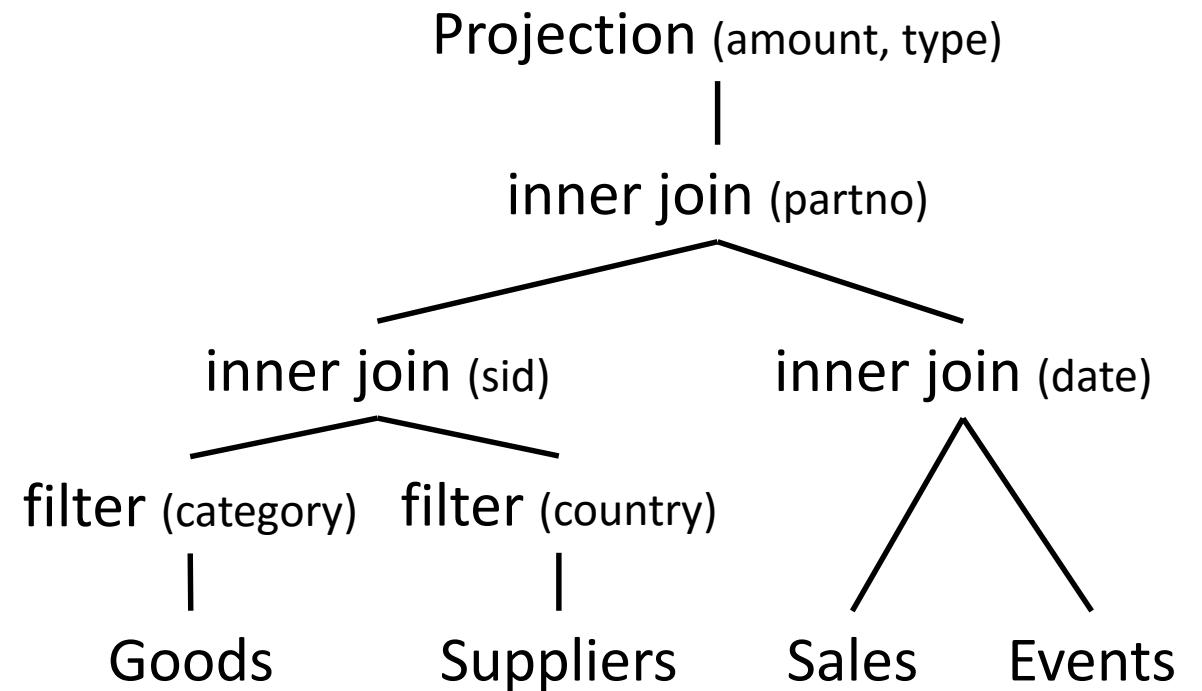
# Understanding current performance

What is the bottleneck? Why?

- Query plan
- Statistics – DBCC SHOW\_STATISTICS
- Estimate statistics yourself
- Experiment with timings

# How would you investigate slow query?

```
SELECT Sales.amount, Events.type
FROM Sales, Events, Goods, Suppliers
WHERE Sales.date = Events.date AND
      Sales.partno = Goods.partno AND
      Suppliers.sid = Goods.sid AND
      Goods.category = 'engine' AND
      Suppliers.country = 'US'
```



# Technique summary

Not a prioritized list!

- Improve system resources
- Improve query optimization
- Don't get in optimizer's way
- **Reduce computation**
- **Reduce data usage**
- Denormalization – next topic
- Eliminating constraints & triggers – generally disrecommended

# Improving system resources

Not always feasible, but can improve EVERY query.

- Faster CPU
- More physical memory
  - Reduce need for I/O
- Faster I/O
  - SSDs instead of disks
  - Distribute over multiple drives to increase bandwidth

# Improving query optimization

Help system optimize better.

- Use simpler synthetic keys
- Create indices – CREATE INDEX
- Manage index growth – FILL FACTOR
- Create specialized statistics – CREATE/UPDATE STATISTICS
- Partition or data stripe table to improve I/O
- Check whether benefits outweigh overheads.

# Overriding the query optimizer

Sometimes, the query optimizer gets it wrong.

- Force it to use index: `SELECT ... FROM ... WITH INDEX ...`
- Force join order

However, the query optimizer is usually smarter than you.




Don't get in optimizer's way

# Use WHERE expressions on the raw attribute

```
SELECT ...  
FROM Account  
WHERE YEAR(created_on) = 2016 AND  
      MONTH(created_on) = 1;
```

Blocks usage of both  
index & statistics.



```
SELECT ...  
FROM Account  
WHERE created_on BETWEEN '1/1/2016' AND '2/1/2016';
```

Other functions commonly used like this: ISNULL() and implicit type conversions.

# Don't loop over SQL calls

```
for each person in person_list:  
    INSERT INTO Person VALUES (this person's data)
```

Use one INSERT INTO with a list of values.

Such examples generally stem from not understanding SQL's capabilities.

```
product_list = SELECT * FROM Product
```

```
for each product in product_list:  
    number_product_sold = SELECT COUNT(product_id)  
                           FROM Sales  
                           WHERE product_id = this product's id
```

Use one SELECT with a join, grouping on product\_id.


# Don't use cursors unnecessarily

Many newbie examples just duplicate SELECT's features.

Reduce computation

# Avoid constant-valued subqueries


```
SELECT first_name, last_name  
FROM Person  
WHERE age = (SELECT Max(age)  
             FROM Person);
```



```
SELECT @oldest_age = MAX(age)  
FROM Person;  
  
SELECT first_name, last_name  
FROM Person  
WHERE age = @oldest_age;
```

# Avoid *correlated* subqueries when possible


```
SELECT name,  
       city,  
       (SELECT company_name FROM Company WHERE id = Customer.company_id)  
FROM Customer;
```



```
SELECT name,  
       city,  
       company_name  
FROM Customer  
INNER JOIN Company ON Company.id = Customer.company_id;
```

# Avoid *correlated* subqueries when possible

```
SELECT id, first_name, last_name, salary
FROM Employee e
WHERE EXISTS (SELECT 1
              FROM Orders o
              WHERE e.id = o.sales_rep_id AND
                    o.customer_id = 123);
```




```
SELECT id, first_name, last_name, salary
FROM Employee e
WHERE id IN (SELECT sales_rep_id
            FROM Orders
            WHERE o.customer_id = 123);
```



# Avoid DISTINCT when it's unnecessary

Find departments with employees

```
SELECT DISTINCT Dept.id, Dept.name  
FROM Dept, Employee  
WHERE Dept.id = Employee.dept_id;
```



```
SELECT id, name  
FROM Dept  
WHERE id IN (SELECT dept_id  
             FROM Employee);
```


A poor variation:

```
SELECT id, name  
FROM Dept  
WHERE id IN (SELECT DISTINCT dept_id  
             FROM Employee);
```

Similarly, avoid UNION when  
UNION ALL is sufficient.

# Don't COUNT to check for existence

```
SELECT id, name  
FROM Dept  
WHERE (SELECT COUNT(*)  
       FROM Employee  
       WHERE Dept.id = Employee.dept_id)  
      > 0;
```




```
SELECT id, name  
FROM Dept  
WHERE id IN (SELECT dept_id  
            FROM Employee);
```

# Use CASE to avoid multiple passes

```
SELECT COUNT (*)  
FROM Employee  
WHERE salary < 20000;
```

```
SELECT COUNT (*)  
FROM Employee  
WHERE salary BETWEEN 20000  
AND 50000;
```

```
SELECT COUNT (*)  
FROM Employee  
WHERE salary > 50000;
```



```
SELECT COUNT (CASE WHEN salary < 20000  
THEN 1 ELSE null END),  
COUNT (CASE WHEN salary BETWEEN 2000  
AND 50000  
THEN 1 ELSE null END),  
COUNT (CASE WHEN salary > 50000  
THEN 1 ELSE null END);
```

# Don't ORDER BY unnecessarily

- Don't sort in nested/helper queries. Only sort your final results.
- Watch for this when using a pre-existing query as a helper.

Reducing data usage

# Don't select too much data

```
SELECT first_name, last_name, city  
FROM Person  
WHERE city = 'Houston';
```


```
SELECT *  
FROM Person  
WHERE city = 'Houston';
```

```
SELECT first_name, last_name, city  
FROM Person;
```

... and then application code  
filters the data.

# SELECT less data in subquery

```
SELECT id, first_name, last_name, salary
FROM Employee e
WHERE dept_id = 456 AND
      id IN (SELECT sales_rep_id FROM Orders);
```



Probably sorts and/or indexes subquery.

In this case, a correlated subquery can be better.

- Subquery results in less data.
- Subquery can use indices on both attributes.

```
SELECT id, first_name, last_name, salary
FROM Employee e
WHERE dept_id = 456 AND
      EXISTS (SELECT 1
              FROM Orders
              WHERE e.id = o.sales_rep_id);
```

# Don't update data with identical data

```
UPDATE Person  
SET active = False  
WHERE ... AND      /* no recent transactions, website visits, or customer calls */  
                active = True;
```

Avoids many writes to Person. Reduces any logging of writes.